

# **CS4224/CS5424 Lecture 6**

## **Data Replication**

# Data Replication

**Objectives:** Improve

- System availability
- Performance
- Scalability

**Challenge:** How to keep replicas synchronized

# Data Replication

- Logical data:  $x$
- Physical copies/replicas:  $x_A, x_B, \dots$ ,
  - ▶  $x_i$  denote the replica at Site  $i$
- Replication transparency
  - ▶ Transaction issues read/write operations on  $x$
  - ▶ Replica control protocol maps operations on logical data to operations on physical replicas

# 1SR & Mutual Consistency

- One-copy database = non-replicated database
- An execution is one-copy serializable (1SR) if it has the same effect as a serial execution on a one-copy database
- A replicated database is in a mutually consistency state if all the replicas of its data items have identical values

# One-Copy Serializability: Example

Assume that initial replicas are synchronized:

$$x_A = x_B, y_A = y_B, z_A = z_B$$

Site A	Site B
$R_1(x_A)$	$R_2(x_B)$
$R_1(y_A)$	$R_2(z_B)$
$W_1(z_A)$	$W_2(y_B)$
$W_2(y_A)$	$W_1(z_B)$
$Commit_1$	$Commit_1$
$Commit_2$	$Commit_2$

# One-Copy Serializability: Example

Assume that initial replicas are synchronized:  $x_A = x_B$

Site A	Site B
$R_1(x_A)$	
$W_1(x_A)$	
$Commit_1$	
	$W_1(x_B)$
	$R_2(x_B)$
$R_3(x_A)$	$W_2(x_B)$
	$Commit_2$
$W_2(x_A)$	
$W_3(x_A)$	
$Commit_3$	
	$W_3(x_B)$

# Mutual Consistency

- A replicated database is in a **mutually consistency** state if all the replicas of its data items have identical values
- **Strong mutual consistency**
  - ▶ All copies of a data item have the same value at the end of update  $X_{act}$
- **Weak mutual consistency**
  - ▶ Does not require all copies of a data item to be identical at the end of update  $X_{act}$
  - ▶ aka **eventual consistency**

# Example 1: MC but not 1SR

- Site A = {x}, Site B = {x, y}, Site C = {x, y, z}
- Assume that the initial replicas are synchronized:  $x_A = x_B = x_C$  &  $y_B = y_C$
- Transactions:
  - ▶  $T_1: W_1(x)$
  - ▶  $T_2: R_2(x), W_2(y)$
  - ▶  $T_3: R_3(x), R_3(y), W_3(z)$
- Local schedules
  - ▶  $S_A: W_1(x_A)$
  - ▶  $S_B: W_1(x_B), R_2(x_B), W_2(y_B)$
  - ▶  $S_C: W_2(y_C), R_3(x_C), R_3(y_C), W_3(z_C), W_1(x_C)$



# Example 2: Neither MC nor 1SR

- Site A =  $\{x\}$ , Site B =  $\{x\}$   
where the initial replicas are synchronized
- Transactions
  - ▶  $T_1: R_1(x), W_1(x)$
  - ▶  $T_2: R_2(x), W_2(x)$
- Local schedules:
  - ▶  $S_A: R_1(x_A), W_1(x_A), W_2(x_A)$
  - ▶  $S_B: R_2(x_B), W_2(x_B), W_1(x_B)$

# One-Copy Serializability

- **Replicated data (RD) schedules**: schedules on replicated database
- **One-copy (1C) schedules**: schedules on non-replicated database
- $T_j$  reads  $x$  from  $T_i$  in a RD schedule if
  1. for some copy  $x_A$  of  $x$ ,  $W_i(x_A)$  precedes  $R_j(x_A)$ , and
  2. there is no  $W_k(x_A)$ ,  $k \neq i$ , that occurs between  $W_i(x_A)$  &  $R_j(x_A)$

# One-Copy Serializability (cont.)

- Let  $T$  denote a set of committed transactions
- Let  $S_{RD}$  denote a RD schedule over  $T$
- Let  $S_{1C}$  denote a 1C schedule over  $T$
- $S_{RD}$  is equivalent to  $S_{1C}$  if
  1.  $T_j$  reads  $x$  from  $T_i$  in  $S_{RD}$  iff  $T_j$  reads  $x$  from  $T_i$  in  $S_{1C}$ , and
  2. for each final write  $W_i(x)$  in  $S_{1C}$ ,  $W_i(x_A)$  is a final write in  $S_{RD}$  for some copy  $x_A$  of  $x$
- A replicated data schedule is **one-copy serializable (1SR)** if it is equivalent to a serial one-copy schedule

# Example 3

- Site A =  $\{x, y\}$ , Site B =  $\{x, y\}$   
where the initial replicas are synchronized
- Transactions
  - ▶  $T_1: W_1(x)$
  - ▶  $T_2: R_2(x), R_2(y)$
  - ▶  $T_3: W_3(y)$
  - ▶  $T_4: R_4(x), R_4(y)$
- Local schedules
  - ▶  $S_A: W_1(x_A), R_2(x_A), R_2(y_A)$
  - ▶  $S_B: W_3(y_B), R_4(x_B), R_4(y_B)$

# Example 4

- Site  $A = \{x, y\}$ , Site  $B = \{x, y\}$   
where the initial replicas are synchronized
- Transactions
  - ▶  $T_1: W_1(x)$
  - ▶  $T_2: R_2(x), R_2(y)$
  - ▶  $T_3: W_3(y)$
  - ▶  $T_4: R_4(x), R_4(y)$
- Local schedules
  - ▶  $S_A: W_1(x_A), R_4(x_A), R_2(x_A), R_2(y_A), R_4(y_A)$
  - ▶  $S_B: W_3(y_B)$

# How to send updates to replicas?

- Suppose a Xact  $T$  has updated data at one site.  
How to send  $T$ 's updates to other replicas?
- Replication Methods
  - ▶ **DBMS-level replication**
    - ★ Statement-based replication
    - ★ Write-ahead log (WAL) shipping
  - ▶ **Application-level replication**
- **Statement-based replication**
  - ▶ Forward  $T$ 's update/insert/delete SQL statements to replica sites for execution
  - ▶ Example: VoltDB

# How to send updates to replicas? (cont.)

- **Write-ahead log (WAL) shipping**
  - ▶ Send T's log records to replica sites for synchronization
    - ★ File-based log shipping
    - ★ Record-based log shipping (streaming replication)
  - ▶ Physical/Logical replication - format of shipped log records are logical/physical
  - ▶ **Physical replication**
    - ★ Storage-based specification of updates (e.g. location of modified bytes on disk block)
    - ★ Examples: Oracle, PostgreSQL (before version 10)
  - ▶ **Logical replication**
    - ★ Contains one log record for each new/deleted/updated tuple
    - ★ Examples: Oracle, MySQL, PostgreSQL (version 10 onwards)

# How to send updates to replicas? (cont.)

- **Application-level replication**
  - ▶ Implement using triggers & stored procedures
  - ▶ More flexibility but higher overhead



# Replication Protocols

		WHERE	
		Centralized	Distributed
WHEN	Eager	Eager Centralized	Eager Distributed
	Lazy	Lazy Centralized	Lazy Distributed

- **When** are updates propagated to copies?
- **Where** are updates allowed to occur?

# When are updates propagated to copies?

- **Eager (or synchronous) update:** Propagates updates to all replicas within context of Xact
- **Lazy (or asynchronous) update:** Xact updates only one replica; updates to remaining replicas are propagated asynchronously

<p><i>T</i>:    Begin transaction       ...       Write(<math>x_a</math>)       Write(<math>x_b</math>)       Write(<math>x_c</math>)       ...       Commit</p>
--

**Eager update**

<p><i>T</i>:    Begin transaction       ...       Write(<math>x_a</math>)       ...       Commit</p>
<p>Sometime later:       Write(<math>x_b</math>)       Write(<math>x_c</math>)</p>

**Lazy update**

# When are updates propagated to copies? (cont.)

- **Eager Update**

- ▶ Enforces strong mutual consistency
- ▶ Based on **Read-One-Write-All (ROWA)** protocols

- **Lazy Update**

- ▶ Xact commits as soon as one replica is updated
- ▶ Updates to remaining replicas are propagated asynchronously
  - ★ **Refresh Xacts** sent to other replica sites after update  
Xact commits
- ▶ Lazy updates from different Xacts can conflict
- ▶ Need to ensure that updates are applied in the **same order** to all replicas

# Where are updates allowed to occur?

- **Centralized techniques**

- ▶ Update is applied to a **master copy** first before propagating to other **slave copies**
  - ★ **Master site**: site that hosts the master copy
  - ★ **Slave site**: site that hosts a slave copy
- ▶ aka **single master**, **master-slave**, or **active-passive** replication

- **Distributed techniques**

- ▶ Update is applied to any copy & then propagated to other copies
- ▶ aka **multimaster**, **update anywhere**, **master-master**, or **active-active** replication

# Assumptions for protocol discussions

- Strict 2PL is used for concurrency control
- Statement-based replication method is used to propagate updates

# Replication Protocols

## 1. Eager Centralized Protocols

- ▶ Eager Single-Master
- ▶ Eager Primary Copy

## 2. Eager Distributed Protocols

## 3. Lazy Centralized Protocols

- ▶ Lazy Single-Master

## 4. Lazy Distributed Protocols

# Eager Single-Master Protocol

- **Eager** = all replicas updated within context of Xact
- **Centralized** = master copy is updated before slave copies

# Eager Single-Master Protocol (cont.)

- There is a **single master site** containing master copies of all objects
  - ▶ TM at master site is the **centralized lock manager**
  - ▶ Coordinating TM for  $T_i$  sends each operation of  $T_i$  to master site's TM
- For each **update operation**  $W_i(O)$ ,
  - ▶ Master copy of  $O$  at master site must be updated first
  - ▶ Updates are then propagated to other copies of  $O$  at slave sites
- For each **read operation**  $R_i(O)$ ,
  - ▶ Coordinating TM for  $T_i$  can read from any one replica of  $O$




# Eager Single-Master Protocol (cont.)

- Site  $S_A$  is the master site
- Consider a Xact  $T_i$  issued at site  $S_B$ 
  - ▶  $TM_B$  is the coordinating TM for  $T_i$
- To process  $R_i(x)$ ,  $TM_B$  sends lock request for  $R_i(x)$  to  $TM_A$
- $TM_A$  checks if S-lock on  $x$  can be granted to  $T_i$
- If granted,
  - ▶  $TM_A$  notifies  $TM_B$  that lock request is granted
  - ▶ If  $TM_B$  has a replica of  $x$ ,  $TM_B$  reads its local copy of  $x$  (i.e.,  $R_i(x_B)$ );
  - ▶ Otherwise,  $TM_B$  sends  $R_i(x)$  to any site (say  $S_C$ ) with a replica of  $x$ 
    - ★  $TM_C$  executes  $R_i(x_C)$  and returns  $x_C$  to  $TM_B$
- Otherwise,  $T_i$  is blocked

# Eager Single-Master Protocol (cont.)

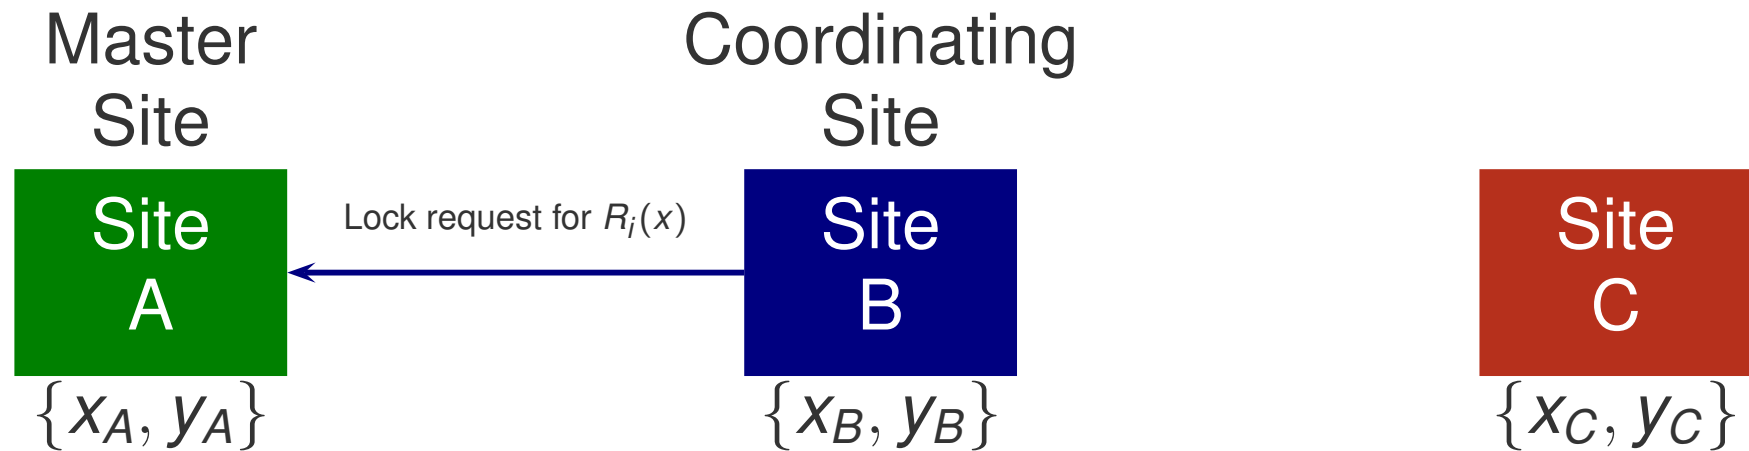
- To process  $W_i(x)$ ,  $TM_B$  sends lock request for  $W_i(x)$  to  $TM_A$
- $TM_A$  checks if X-lock on  $x$  can be granted to  $T_i$
- If granted,
  - ▶  $TM_A$  updates its copy of  $x$  (i.e.,  $W_i(x_A)$ )
  - ▶  $TM_A$  notifies  $TM_B$  that lock request is granted
  - ▶  $TM_B$  sends  $W_i(x)$  to other sites with replicas of  $x$
  - ▶  $TM_B$  executes  $W_i(x_B)$  if  $S_B$  has a replica of  $x$
- Otherwise,  $T_i$  is blocked

# Eager Single-Master Protocol: Example



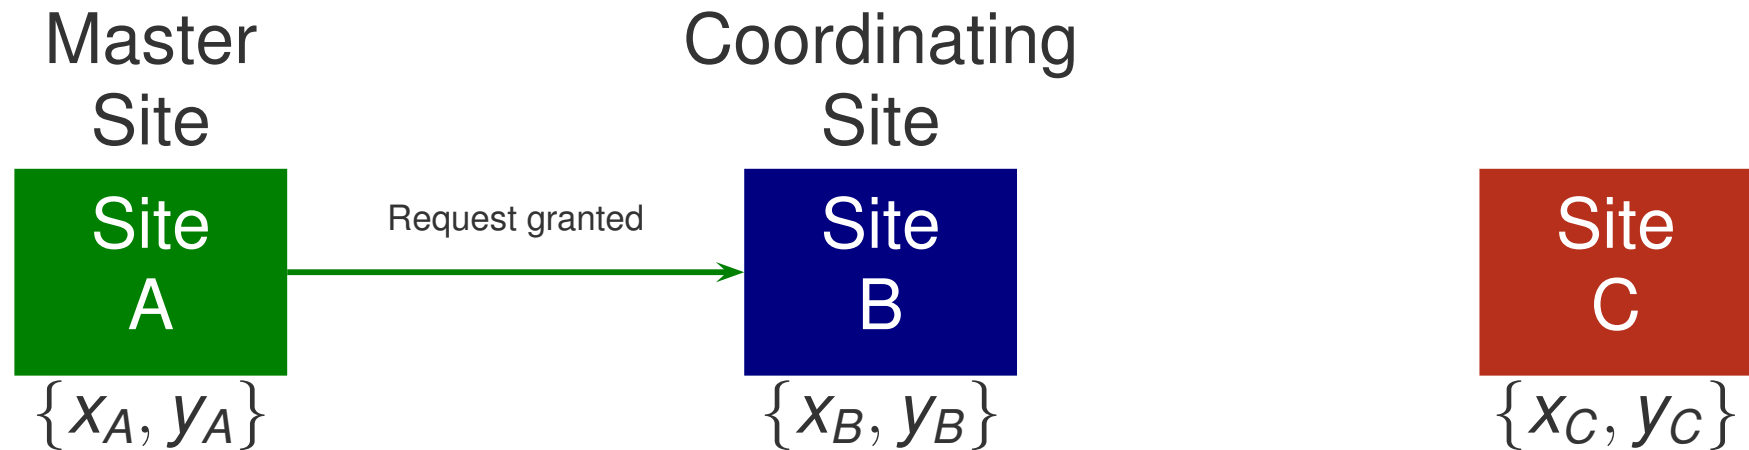
Transaction  $T_i$  needs to read  $x$  & write  $y$   
**Site B** is the coordinating site for Xact  $T_i$

# Eager Single-Master Protocol: Example



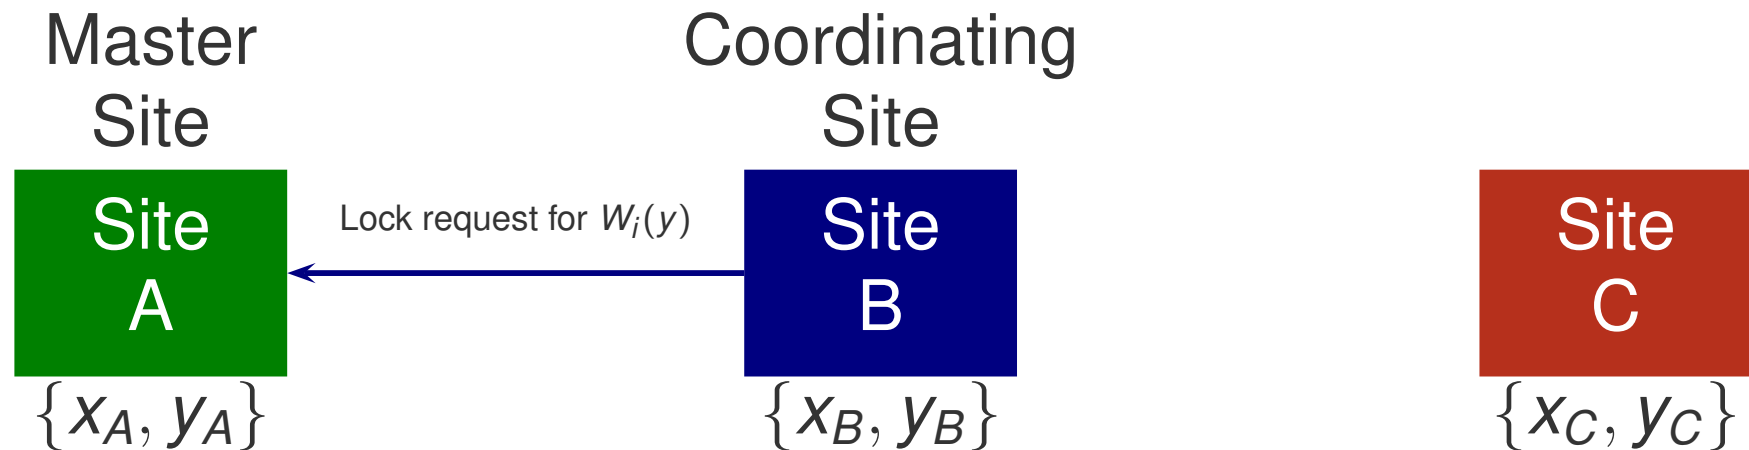
$TM_B$  sends lock request for  $R_i(x)$  to  $TM_A$

# Eager Single-Master Protocol: Example




$TM_A$  grants S-lock for  $R_i(x)$  & notifies  $TM_B$

# Eager Single-Master Protocol: Example



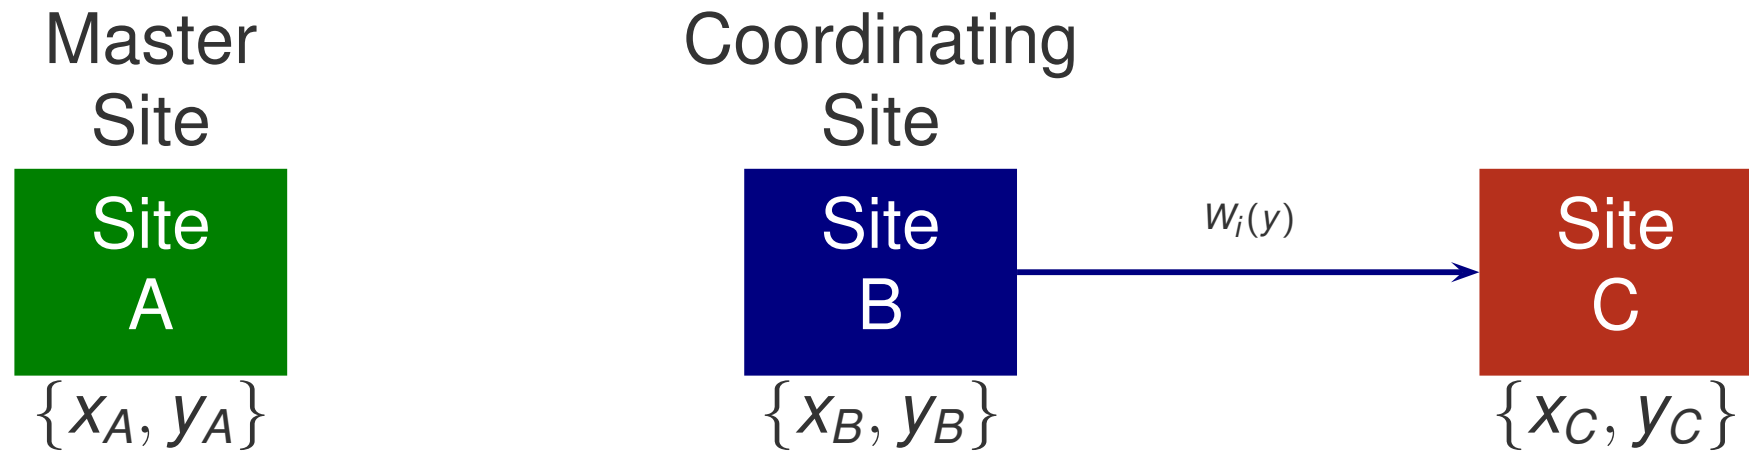
$TM_B$  executes  $R_i(x_B)$  &  
 $TM_B$  sends lock request for  $W_i(y)$  to  $TM_A$

# Eager Single-Master Protocol: Example



$TM_A$  grants X-lock for  $W_i(y)$ , executes  $W_i(y_A)$  & notifies  $TM_B$


# Eager Single-Master Protocol: Example



$TM_B$  executes  $W_i(y_B)$  & sends  $W_i(y)$  to  $TM_C$



# Eager Single-Master Protocol: Example



$TM_C$  executes  $W_i(y_C)$

# Eager Primary Copy Protocol

- Generalization of Eager Single-Master Protocol
- For each replicated object, one of its copies is designated the **primary copy**
- The master site for object  $O$  is the site that stores the primary copy of  $O$
- Each master site runs a lock manager
  - ▶ Controls lock requests/releases for primary copies under its control

# Eager Primary Copy Protocol: Example

Master Site  
for x



$\{x_A, y_A\}$

Coordinating  
Site



$\{x_B, y_B\}$

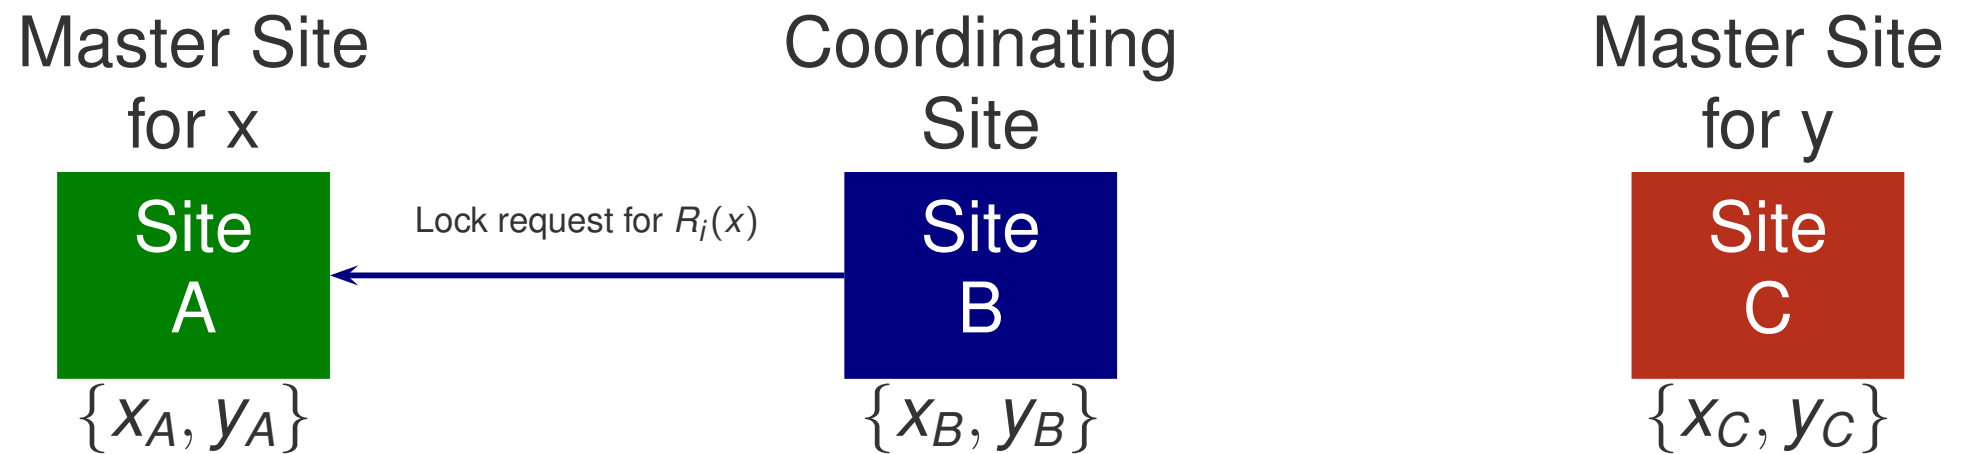
Master Site  
for y



$\{x_C, y_C\}$


Transaction  $T_i$  needs to read  $x$  & write  $y$   
**Site B** is the coordinating site for Xact  $T_i$

# Eager Primary Copy Protocol: Example



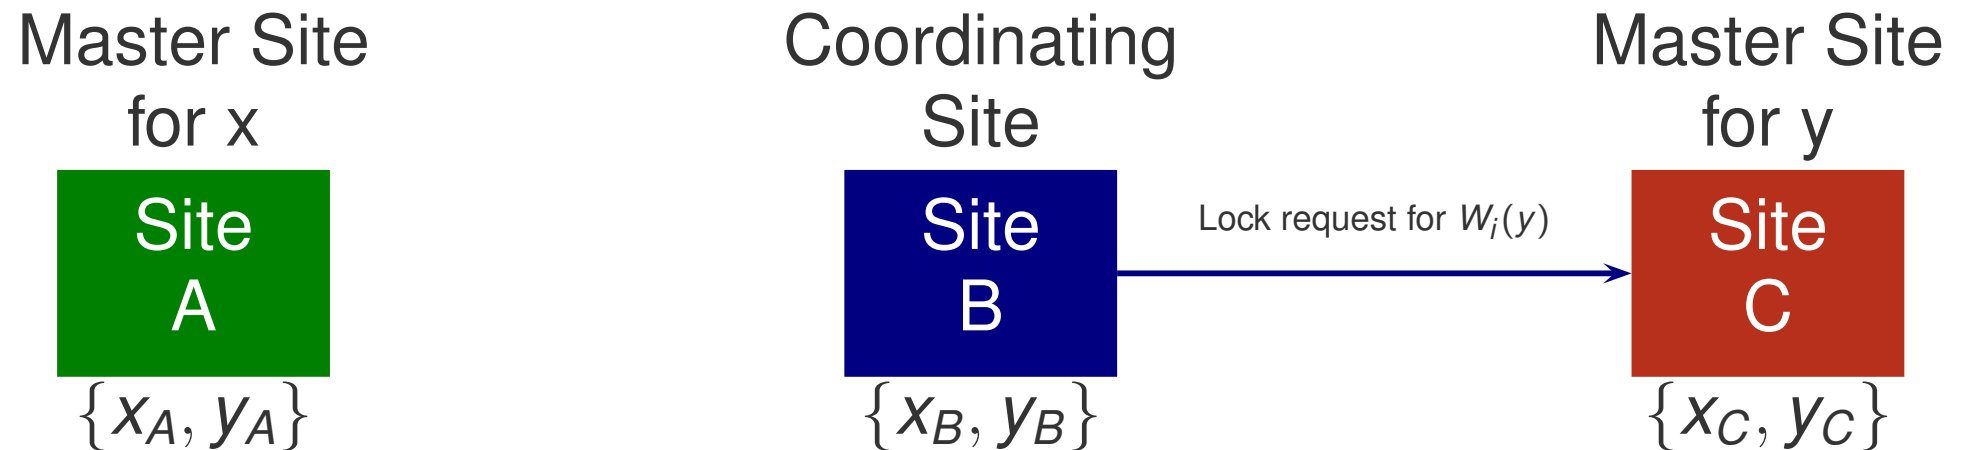
$TM_B$  sends lock request for  $R_i(x)$  to  $TM_A$

# Eager Primary Copy Protocol: Example




$TM_A$  grants S-lock for  $R_i(x)$  & notifies  $TM_B$

# Eager Primary Copy Protocol: Example




$TM_B$  executes  $R_i(x_B)$  &  
sends lock request for  $W_i(y)$  to  $TM_C$

# Eager Primary Copy Protocol: Example



$TM_C$  grants X-lock for  $W_i(y)$ , executes  $W_i(y_C)$  & notifies  $TM_B$

# Eager Primary Copy Protocol: Example



$TM_B$  executes  $W_i(y_B)$  & sends  $W_i(y)$  to  $TM_A$



# Eager Primary Copy Protocol: Example

Master Site  
for x

Site  
A

$\{x_A, y_A\}$

Coordinating  
Site

Site  
B

$\{x_B, y_B\}$

Master Site  
for y

Site  
C

$\{x_C, y_C\}$

$TM_A$  executes  $W_i(y_A)$

# Eager Distributed Protocols

- **Eager** = all replicas updated within context of Xact
- **Distributed** = any replica can be updated first
- Each site runs a lock manager
  - ▶ Controls lock requests/releases for its local replicas


# Eager Distributed Protocols (cont.)

- Consider a Xact  $T_i$  issued at site  $S_B$ 
  - ▶  $TM_B$  is the coordinating TM for  $T_i$
- To process  $R_i(x)$ ,
  - ▶ If site  $S_B$  has a local replica of  $x$ 
    - ★  $TM_B$  checks if S-lock on  $x$  can be granted for  $T_i$
    - ★ If S-lock for  $R_i(x)$  is granted,  $TM_B$  reads from its local replica of  $x$  (i.e.,  $R_i(x_B)$ )
    - ★ Otherwise,  $T_i$  is blocked
  - ▶ Else
    - ★  $TM_B$  sends  $R_i(x)$  to any site (say  $S_C$ ) with a copy of  $x$
    - ★ If S-lock for  $R_i(x)$  is granted by  $TM_C$ ,  $TM_C$  reads  $x_C$  & returns  $x_C$  to  $TM_B$
    - ★ Otherwise,  $T_i$  is blocked

# Eager Distributed Protocols (cont.)


- To process  $W_i(x)$ ,
  - ▶ If site  $S_B$  has a local replica of  $x$ 
    - ★  $TM_B$  checks if X-lock on  $x$  can be granted for  $T_i$
    - ★ If X-lock for  $W_i(x)$  is granted,  $TM_B$  updates its local replica of  $x$  (i.e.,  $W_i(x_B)$ ) & sends  $W_i(x)$  to other sites with replicas of  $x$
    - ★ Otherwise,  $T_i$  is blocked
  - ▶ Else
    - ★  $TM_B$  sends  $W_i(x)$  to any site (say  $S_C$ ) with a copy of  $x$
    - ★ If X-lock for  $W_i(x)$  is granted by  $TM_C$ ,  $TM_C$  updates  $x_C$  & notifies  $TM_B$  that request is granted.  $TM_B$  sends  $W_i(x)$  to other sites with replicas of  $x$
    - ★ Otherwise,  $T_i$  is blocked

# Eager Distributed Protocol: Example



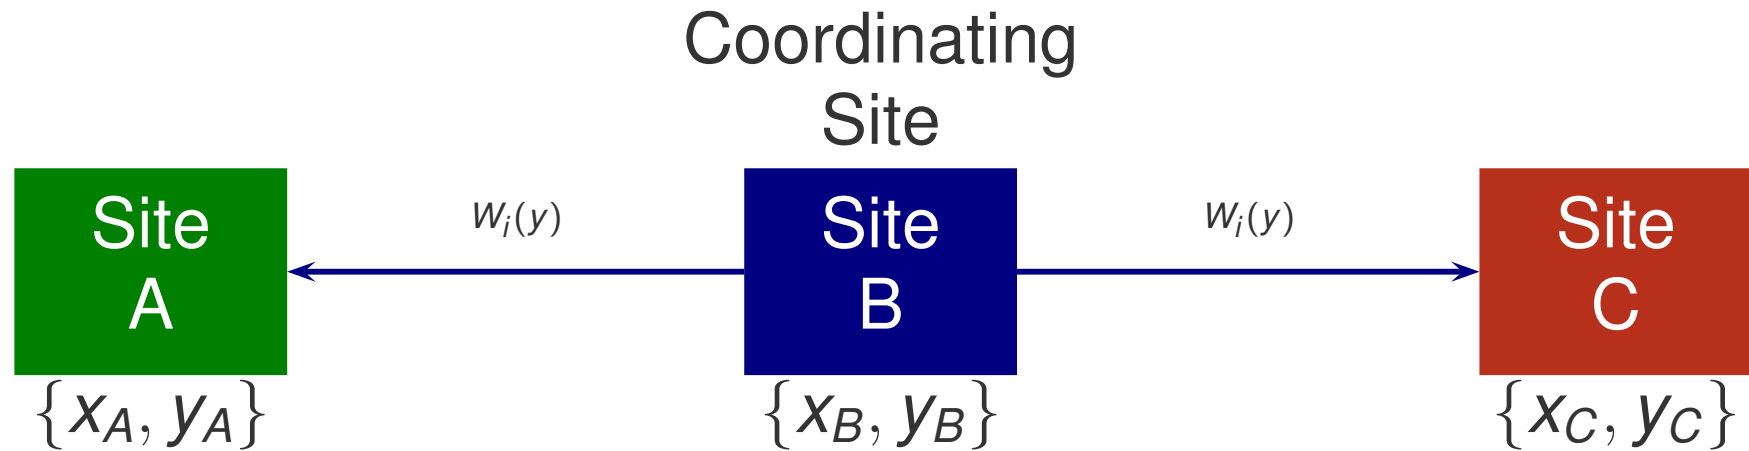
Transaction  $T_i$  needs to read  $x$  & write  $y$   
**Site B** is the coordinating site for Xact  $T_i$

# Eager Distributed Protocol: Example




For  $R_i(x)$ ,  $TM_B$  grants S-lock for  $R_i(x)$  & executes  $R_i(x_B)$

# Eager Distributed Protocol: Example



For  $W_i(y)$ ,  $TM_B$  grants X-lock for  $W_i(y)$ , executes  $W_i(y_B)$  & sends  $W_i(y)$  to  $TM_A$  &  $TM_C$

# Eager Distributed Protocol: Example



$TM_A$  grants X-lock for  $W_i(y)$  & executes  $W_i(y_A)$

$TM_C$  grants X-lock for  $W_i(y)$  & executes  $W_i(y_C)$



# Lazy Centralized Protocols

- **Lazy** = Only one replica is updated within context of Xact; other replicas are updated asynchronously
  - ▶ Refresh Xacts sent to other replica sites after update Xact commits
- **Centralized** = master copy is updated before slave copies

# Lazy Single-Master Protocol

- There is a **single master site** containing master copies of all objects
- For each **update operation**  $W_i(O)$ ,
  - ▶ Master copy of  $O$  at master site must be updated first
  - ▶ Updates are then propagated asynchronously to other copies of  $O$  after  $T_i$  commits
- For each **read operation**  $R_i(O)$ ,
  - ▶ Coordinating TM for  $T_i$  can read from any one replica of  $O$

# Lazy Single-Master Protocol (cont.)

- Site  $S_A$  is the master site
- Consider a Xact  $T_i$  issued at site  $S_B$ 
  - ▶  $TM_B$  is the coordinating TM for  $T_i$
- To process  $R_i(x)$ ,  $TM_B$  sends lock request for  $R_i(x)$  to  $TM_A$
- $TM_A$  checks if S-lock on  $x$  can be granted to  $T_i$
- If granted,
  - ▶  $TM_A$  notifies  $TM_B$  that lock request is granted
  - ▶ If  $TM_B$  has a replica of  $x$ ,  $TM_B$  reads its local copy of  $x$  (i.e.,  $R_i(x_B)$ );
  - ▶ Otherwise,  $TM_B$  sends  $R_i(x)$  to any site (say  $S_C$ ) with a replica of  $x$ 
    - ★  $TM_C$  executes  $R_i(x_C)$  & returns  $x_C$  to  $TM_B$
- Otherwise,  $T_i$  is blocked

# Lazy Single-Master Protocol (cont.)

- To process  $W_i(x)$ ,  $TM_B$  sends lock request for  $W_i(x)$  to  $TM_A$
- $TM_A$  checks if X-lock on  $x$  can be granted to  $T_i$
- If granted,
  - ▶  $TM_A$  updates its copy of  $x$  (i.e.,  $W_i(x_A)$ )
  - ▶  $TM_A$  notifies  $TM_B$  that lock request is granted
- Otherwise,  $T_i$  is blocked

# Lazy Single-Master Protocol (cont.)

- When  $T_i$  commits,  $TM_B$  sends  $Commit_i$  to  $TM_A$
- $TM_A$  executes  $Commit_i$  & releases locks for  $T_i$
- $TM_A$  checks if X-locks can be granted for  $T_i$ 's refresh transactions
- If granted,  $TM_A$  sends refresh transactions to other sites to propagate  $T_i$ 's updates; otherwise, the sending of refresh transactions is blocked

# Refresh Transactions


- Let  $T_i^r$  denote the refresh transaction for  $T_i$
- Each  $T_i^r$  is executed as a local refresh transaction at a slave site
- Important for refresh Xacts to be applied to all slave sites in the same order
  - ▶ If  $TS(T_i^r)$  &  $TS(T_j^r)$  are both required to be executed at multiple sites, then they must be executed in the same order at these sites
- Slave site applies refresh transactions in timestamp order as follows
  - ▶ Each  $T_i^r$  has a timestamp denoted by  $TS(T_i^r)$
  - ▶  $TS(T_i^r) = commitTS(T_i)$  which is the commit timestamp of  $T_i$ 
    - ★  $T_i$  commits before  $T_j$  iff  $commitTS(T_i) < commitTS(T_j)$
  - ▶  $T_i^r$  is executed before  $T_j^r$  at a site iff  $TS(T_i^r) < TS(T_j^r)$

# Refresh Transactions: Example

- DDBMS: Site A =  $\{x, y\}$ , Site B =  $\{x\}$ , Site C =  $\{x, y\}$
- Site A is the master site
- $TM_B$  is the coordinator for the execution of  $T_i$ 
  - ▶  $T_i: R_i(x), R_i(y), W_i(x), W_i(y)$
  - ▶  $T_i^r$  consists of the set of updates  $\{W_i^r(x), W_i^r(y)\}$
- Local Schedules:
 

$S_A:$	$W_i(x_A), W_i(y_A), C_i$	
$S_B:$	$R_i(x_B),$	$W_i^r(x_B), C_i^r$
$S_C:$	$R_i(y_C),$	$W_i^r(x_C), W_i^r(y_C), C_i^r$
- Note:  $C_i$  &  $C_i^r$  denote the commit of  $T_i$  &  $T_i^r$ , respectively

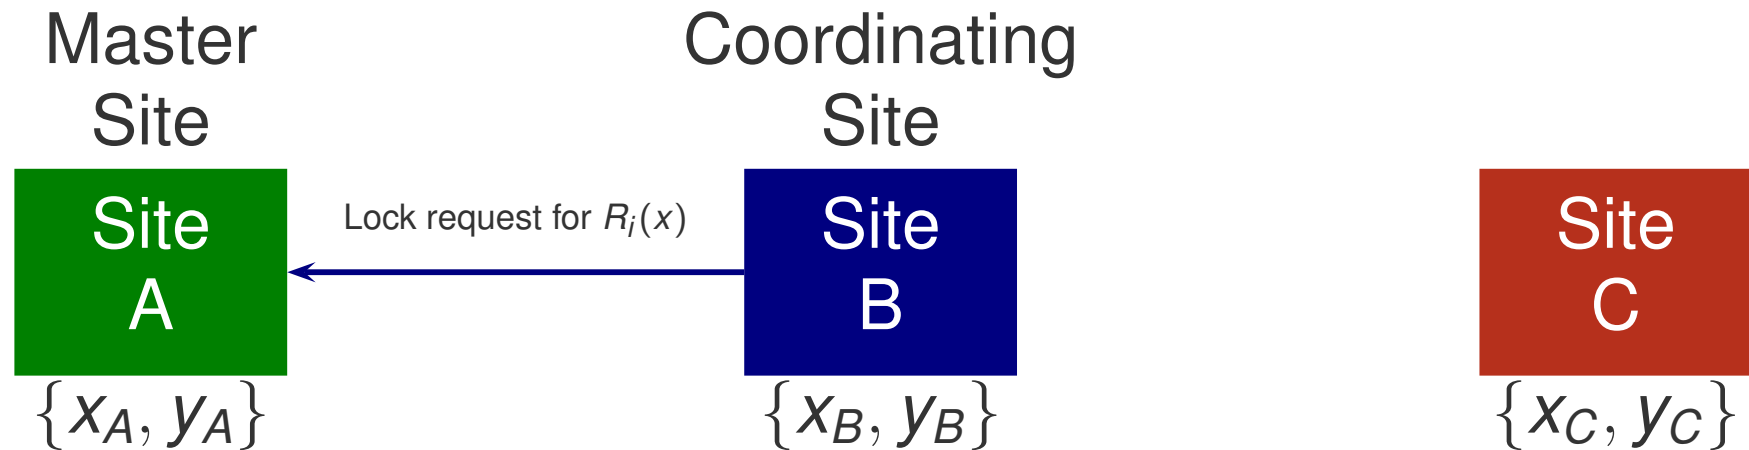
# Lazy Single-Master Protocol: Example



Transaction  $T_i$  needs to read  $x$  & write  $y$   
**Site B** is the coordinating site for Xact  $T_i$




# Lazy Single-Master Protocol: Example



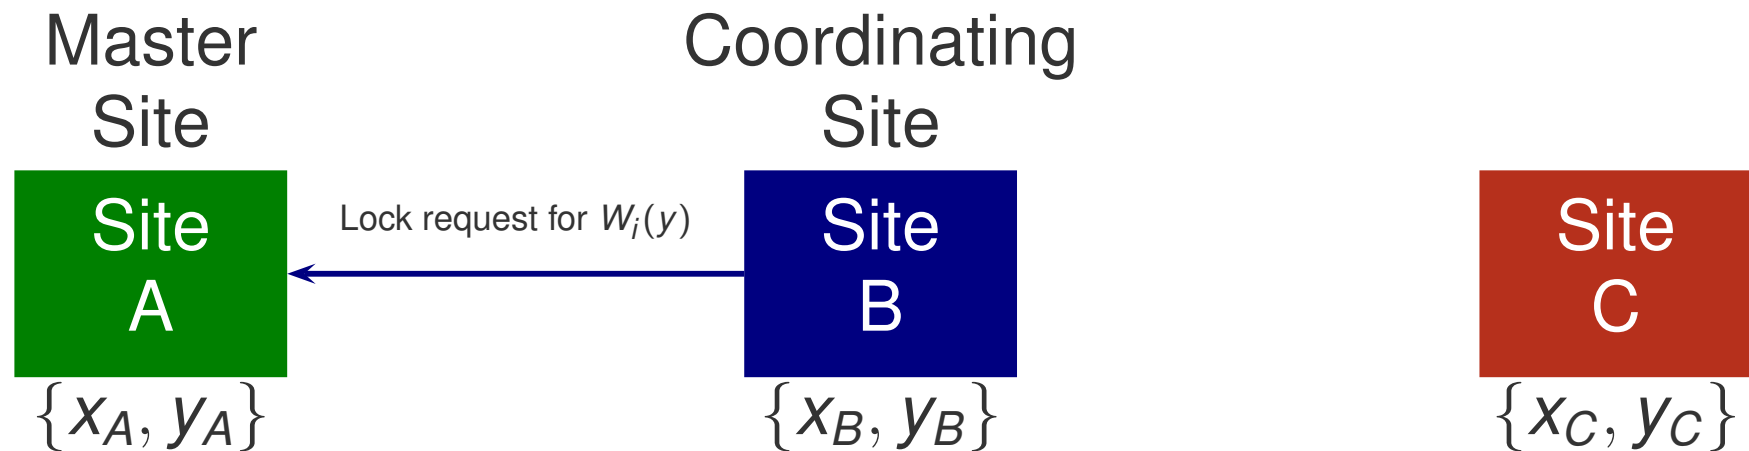
$TM_B$  sends lock request for  $R_i(x)$  to  $TM_A$

# Lazy Single-Master Protocol: Example



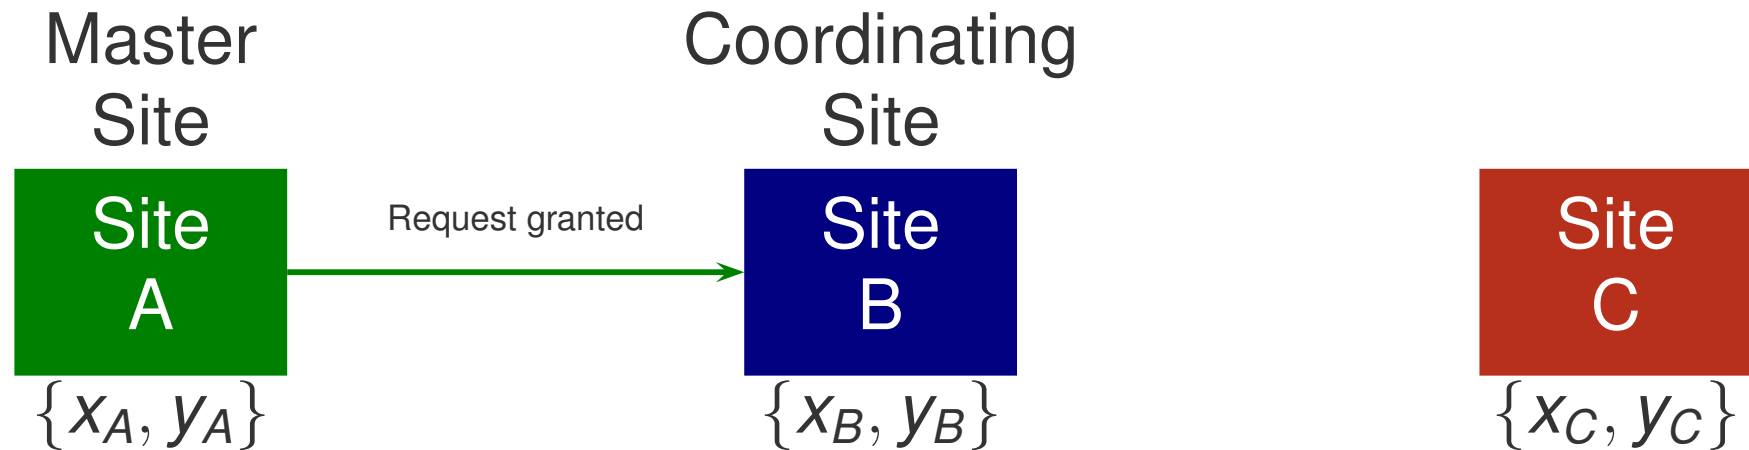
$TM_A$  grants S-lock for  $R_i(x)$  & notifies  $TM_B$

# Lazy Single-Master Protocol: Example



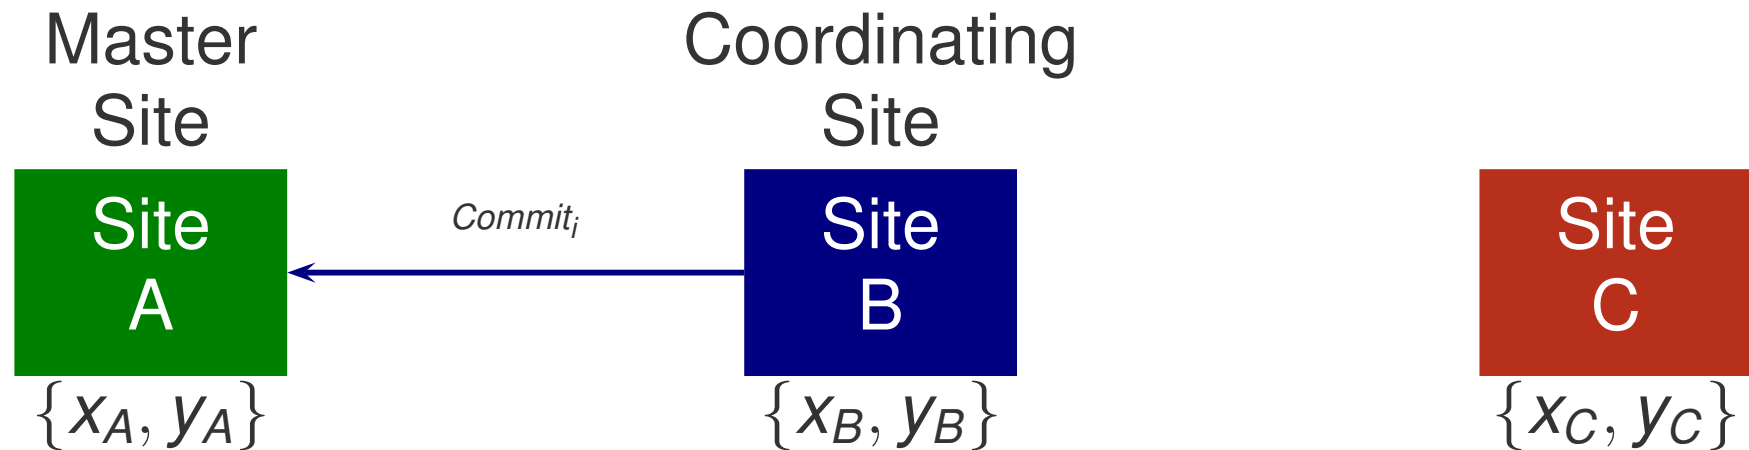
$TM_B$  executes  $R_i(x_B)$  & sends lock request for  $W_i(y)$  to  $TM_A$

# Lazy Single-Master Protocol: Example



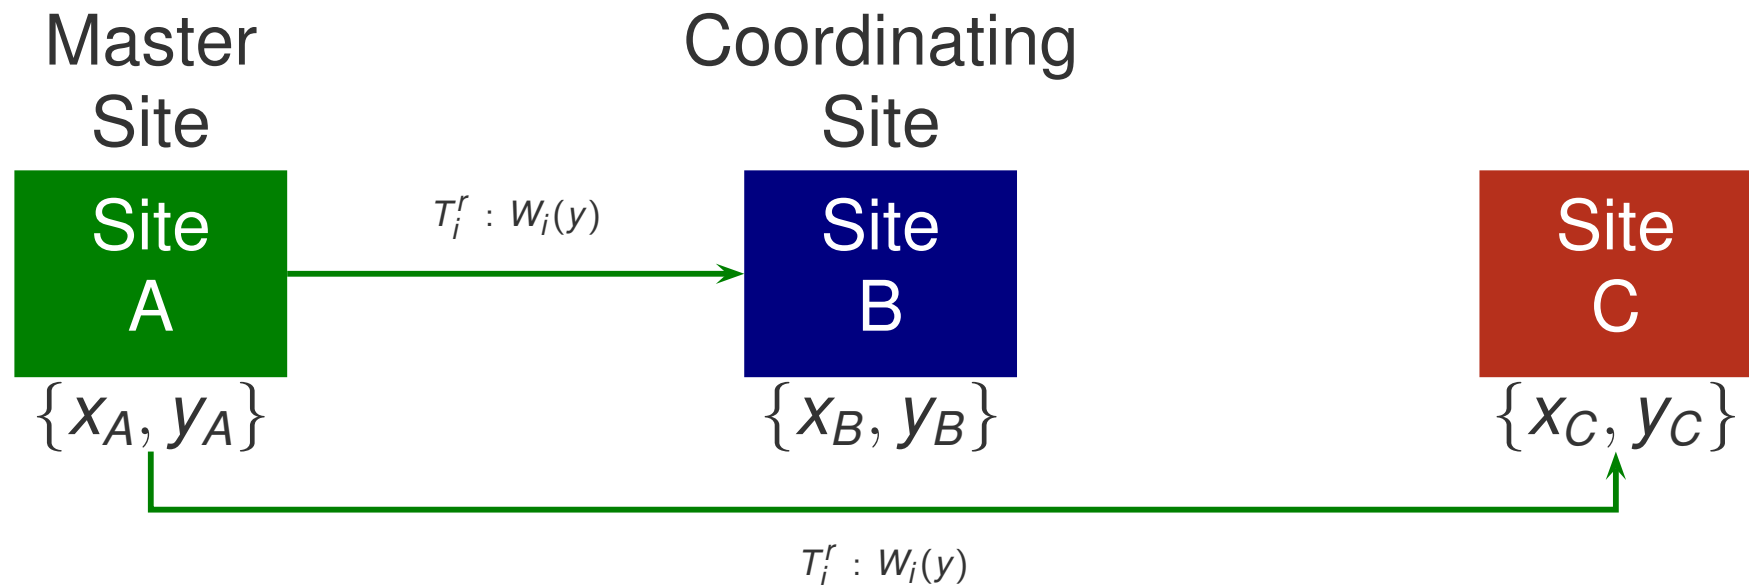
$TM_A$  grants X-lock for  $W_i(y)$ , executes  $W_i(y_A)$  & notifies  $TM_B$

# Lazy Single-Master Protocol: Example




$TM_B$  sends  $Commit_i$  to  $TM_A$

# Lazy Single-Master Protocol: Example



$TM_A$  executes  $Commit_i$ , releases locks for  $T_i$ , grants X-lock for  $W_i^r(y)$ , & sends  $T_i^r : W_i(y)$  to  $TM_B$  &  $TM_C$

# Lazy Single-Master Protocol: Example



$TM_B$  executes  $W_i^r(y_B)$   
 $TM_C$  executes  $W_i^r(y_C)$

# Example 1

- Site A contains master copies of  $\{x, y\}$
- Site B contains slave copies of  $\{x, y\}$
- Transactions
  - ▶  $T_1: R_1(x), W_1(y)$
  - ▶  $T_2: W_2(x), W_2(y)$
- Assume  $T_1$  issued at Site A &  $T_2$  issued at Site B
- Local schedules:  
 $S_A: R_1(x_A), W_1(y_A), C_1, W_2(x_A), W_2(y_A), C_2$   
 $S_B: W_1^r(y_B), C_1^r, W_2^r(x_B), W_2^r(y_B), C_2^r$
- Note that  $TS(T_1^r) < TS(T_2^r)$



# Example 2

- Site A contains master copies of  $\{x, y\}$
- Site B contains slave copies of  $\{x, y\}$
- Transactions
  - ▶  $T_1: R_1(x), W_1(y)$
  - ▶  $T_2: W_2(x), W_2(y)$
- Assume  $T_1$  &  $T_2$  are issued at Site B
- Local schedules:

$S_A : W_2(x_A), W_2(y_A), C_2,$   
 $S_B : \quad \quad \quad R_1(x_B), \quad \quad \quad W_1(y_A), \quad C_1$   
 $\quad \quad \quad \quad \quad \quad \quad \quad \quad \quad \quad W_2^r(x_B), W_2^r(y_B), C_2^r, W_1^r(y_B), C_1^r$

- Note that  $TS(T_2^r) < TS(T_1^r)$

# Example 3

- Site A contains master copy of  $x$
- Site B contains slave copy of  $x$
- Transactions:
  - ▶  $T_1: W_1(x), R_1(x)$
- Assume  $T_1$  issued at Site B
- Local schedules:

$$\begin{array}{l} S_A : W_1(x_A), \\ S_B : R_1(x_B), \quad C_1 W_1^r(x_B), C_1^r \end{array}$$

# Lazy Distributed Protocols

- **Lazy** = Only one replica is updated within context of Xact; other replicas are updated asynchronously
- **Distributed** = any replica can be updated first
- Each site runs a lock manager
  - ▶ Controls lock requests/releases for its local replicas


# Lazy Distributed Protocols (cont.)

- Consider a Xact  $T_i$  issued at site  $S_B$ 
  - ▶  $TM_B$  is the coordinating TM for  $T_i$
- To process  $R_i(x)$ ,
  - ▶ If site  $S_B$  has a local replica of  $x$ 
    - ★  $TM_B$  checks if S-lock on  $x$  can be granted for  $T_i$
    - ★ If S-lock for  $R_i(x)$  is granted,  $TM_B$  reads from its local replica of  $x$  (i.e.,  $R_i(x_B)$ )
    - ★ Otherwise,  $T_i$  is blocked
  - ▶ Else
    - ★  $TM_B$  sends  $R_i(x)$  to any site (say  $S_C$ ) with a copy of  $x$
    - ★ If S-lock for  $R_i(x)$  is granted by  $TM_C$ ,  $TM_C$  executes  $R_i(x_C)$  & returns  $x_C$  to  $TM_B$
    - ★ Otherwise,  $T_i$  is blocked

# Lazy Distributed Protocols (cont.)


- To process  $W_i(x)$ ,
  - ▶ If site  $S_B$  has a local replica of  $x$ 
    - ★  $TM_B$  checks if X-lock on  $x$  can be granted for  $T_i$
    - ★ If X-lock for  $W_i(x)$  is granted,  $TM_B$  updates its local replica of  $x$  (i.e.,  $W_i(x_B)$ )
    - ★ Otherwise,  $T_i$  is blocked
  - ▶ Else
    - ★  $TM_B$  sends  $W_i(x)$  to any site (say  $S_C$ ) with a copy of  $x$
    - ★ If X-lock for  $W_i(x)$  is granted by  $TM_C$ ,  $TM_C$  updates  $x_C$  (i.e.,  $W_i(x_C)$ )
    - ★ Otherwise,  $T_i$  is blocked
- When  $T_i$  commits,
  - ▶  $TM_B$  executes  $Commit_i$ , releases locks for  $T_i$ , & sends **refresh transactions** to other sites to propagate  $T_i$ 's updates

# Lazy Distributed Protocol: Example




Transaction  $T_i$  needs to read  $x$  & write  $y$   
**Site B** is the coordinating site for Xact  $T_i$

# Lazy Distributed Protocol: Example



For  $R_i(x)$ ,  $TM_B$  grants S-lock for  $R_i(x)$  & executes  $R_i(x_B)$


# Lazy Distributed Protocol: Example



For  $W_i(y)$ ,  $TM_B$  grants X-lock for  $W_i(y)$  & executes  $W_i(y_B)$




# Lazy Distributed Protocol: Example



$TM_B$  executes  $Commit_i$ , releases locks for  $T_i$  & sends  $T_i^r: W_i(y)$  to  $TM_A$  &  $TM_C$

# Lazy Distributed Protocol: Example



$TM_A$  grants X-lock for  $W_i^r(y)$  & executes  $W_i^r(y_A)$   
 $TM_C$  grants X-lock for  $W_i^r(y)$  & executes  $W_i^r(y_C)$

# Reconciliation of Inconsistent Updates

- Multiple Xacts could update different copies of same data concurrently at different sites
  - ▶ Conflicting updates can occur!

- **Example:**

$$\begin{array}{lcl} S_A: & W_1(x_A), & C_1, \quad W_2^r(x_A), C_2^r \\ S_B: & W_2(x_B), & C_2, \quad W_1^r(x_B), C_1^r \end{array}$$

- Requires reconciliation procedure
  - ▶ **Last-Writer-Wins heuristic** (a.k.a. timestamp order heuristic):  
apply updates in timestamp order

# Last-Writer-Wins Heuristic

- Used to reconcile inconsistent updates
- **Last-Writer-Wins Heuristic**
  - ▶ If there are two concurrent updates  $W_i(x)$  &  $W_j(x)$ ,  $W_i(x)$  wins if  $TS(T_j^r) < TS(T_i^r)$ 
    - ★  $W_j(x)$  is ignored if  $x$  was last updated by  $T_i$  &  $TS(T_j^r) < TS(T_i^r)$
- **Example:**  $x_A = x_B = 1$ ,  $TS(T_1^r) < TS(T_2^r)$

Site A	Site B	Comments
$W_1(x_A, 10)$		$x_A = 10$
	$W_2(x_B, 20)$	$x_B = 20$
$Commit_1$	$Commit_2$	
	Receives $W_1^r(x_B, 10, TS(T_1^r))$	Xact is ignored
Receives $W_2^r(x_A, 20, TS(T_2^r))$		
$W_2^r(x_A, 20, TS(T_2^r))$ $Commit_2^r$		$x_A = 20$

# Last-Writer-Wins Heuristic (cont.)

- Heuristic only works for updates that are blind writes
- An update  $W(x)$  is a **blind write** if the new value of  $x$  is computed independent of its previous value



# Example: Non-Blind Updates

- $T_1: R_1(x), x = x \times 100, W_1(x)$
- $T_2: R_2(x), x = x + 10, W_2(x)$
- $x_A = 1, x_B = 1, TS(T_1^r) < TS(T_2^r)$



Site A	Site B	Comments
$R_1(x_A), W_1(x_A, 100)$		$x_A = 100$
	$R_2(x_B), W_2(x_B, 11)$	$x_B = 11$
$Commit_1$	$Commit_2$	
	Receives $W_1^r(x_B, 100, TS(T_1^r))$	Xact is ignored
Receives $W_2^r(x_A, 11, TS(T_2^r))$		
$W_2^r(x_A, 11, TS(T_2^r))$ $Commit_2^r$		$x_A = 11$

# Quorum Consensus (QC) Protocol

ROWA  
Protocol



QC  
Protocol



# Quorum Consensus (QC) Protocol

- Each copy  $O_i$  of an object  $O$  is assigned a non-negative weight,  $Wt(O_i)$
- $Wt(O)$  = total weight of all copies of  $O$   
$$Wt(O) = \sum_{O_i \in S} Wt(O_i), \quad S = \text{set of all copies of } O$$
- Each object  $O$  is associated with two thresholds:
  - ▶ Read threshold  $T_r(O)$
  - ▶ Write threshold  $T_w(O)$



# Quorum Consensus (QC) Protocol (cont.)

- A **read quorum** for object  $O$ ,  $Q_r(O)$ , is a subset of copies of  $O$  such that

$$\sum_{O_i \in Q_r(O)} Wt(O_i) \geq T_r(O)$$

- A **write quorum** for object  $O$ ,  $Q_w(O)$ , is a subset of copies of  $O$  such that

$$\sum_{O_i \in Q_w(O)} Wt(O_i) \geq T_w(O)$$

- The read & write thresholds for object  $O$  satisfy the following constraints:
  1.  $T_r(O) + T_w(O) > Wt(O)$
  2.  $2 \times T_w(O) > Wt(O)$

# Quorum Consensus (QC) Protocol (cont.)

- Each object copy is associated with a **version number** to indicate how up-to-date its value is
  - ▶ Higher version number means more up-to-date
  - ▶ Version number is initialized to 0

# Quorum Consensus (QC) Protocol (cont.)

- To read an object  $O$ ,
  - ▶ Acquire S-locks on a read quorum for  $O$ ,  $Q_r(O)$
  - ▶ Read all copies in  $Q_r(O)$  and return the copy with the highest version number
- To write an object  $O$ ,
  - ▶ Acquire X-locks on a write quorum for  $O$ ,  $Q_w(O)$
  - ▶ Let  $n$  be the highest version number among all copies in  $Q_w(O)$
  - ▶ Write all copies in  $Q_w(O)$  and update their version numbers to  $n + 1$

# QC Protocol: Example

Replica	Weight	Value	Version
$X_A$	1	10	0
$X_B$	1	10	0
$X_C$	4	10	0
$X_D$	2	10	0
$X_E$	3	10	0
$X_F$	1	10	0
$X_G$	1	10	0

$$Wt(x) = 13, T_r(x) = 5, T_w(x) = 9$$

# QC Protocol: Example (cont.)

Operation:  $W_1(x, 15)$

Replica	Weight	Value	Version
$X_A$	1	10	0
$X_B$	1	10	0
$X_C$	4	<del>10</del> 15	<del>0</del> 1
$X_D$	2	<del>10</del> 15	<del>0</del> 1
$X_E$	3	<del>10</del> 15	<del>0</del> 1
$X_F$	1	10	0
$X_G$	1	10	0

$$Wt(x) = 13, T_r(x) = 5, T_w(x) = 9$$

# QC Protocol: Example (cont.)

Operation:  $R_2(x)$

Replica	Weight	Value	Version
$x_A$	1	10	0
$x_B$	1	10	0
$x_C$	4	15	1
$x_D$	2	15	1
$x_E$	3	15	1
$x_F$	1	10	0
$x_G$	1	10	0

$$Wt(x) = 13, T_r(x) = 5, T_w(x) = 9$$

# QC Protocol: Example (cont.)

Operation:  $R_3(x)$

Replica	Weight	Value	Version
$X_A$	1	10	0
$X_B$	1	10	0
$X_C$	4	15	1
$X_D$	2	15	1
$X_E$	3	15	1
$X_F$	1	10	0
$X_G$	1	10	0

$$Wt(x) = 13, T_r(x) = 5, T_w(x) = 9$$

# QC Protocol: Example (cont.)

Operation:  $W_3(x, 30)$

Replica	Weight	Value	Version
$X_A$	1	10	0
$X_B$	1	<del>10</del> 30	<del>0</del> 2
$X_C$	4	<del>15</del> 30	<del>1</del> 2
$X_D$	2	<del>15</del> 30	<del>1</del> 2
$X_E$	3	<del>15</del> 30	<del>1</del> 2
$X_F$	1	10	0
$X_G$	1	10	0

$$Wt(x) = 13, T_r(x) = 5, T_w(x) = 9$$



# QC Protocol: Example (cont.)

Operation:  $R_5(x)$

Replica	Weight	Value	Version
$X_A$	1	10	0
$X_B$	1	30	2
$X_C$	4	30	2
$X_D$	2	30	2
$X_E$	3	30	2
$X_F$	1	10	0
$X_G$	1	10	0

$$Wt(x) = 13, T_r(x) = 5, T_w(x) = 9$$

# Handling Failures


- Detect site failures using timeout mechanism
- Lecture focuses on centralized replication protocol
  - ▶ Single-master replication

# Failure of Slave Sites


- Suppose some slave site(s) have failed
- **Lazy replication**
  - ▶ Synchronize unavailable replicas later when they become available
- **Eager replication**
  - ▶ Eager replication techniques are based on ROWA protocol
  - ▶ Drawback of ROWA: Update Xact can't terminate even if one replica is unavailable
  - ▶ **Read-One/Write-All Available (ROWAA) protocol**
    - ★ Relax ROWA protocol to increase availability
    - ★ Update all available replicas & terminate Xact
    - ★ Synchronize unavailable replicas later when they become available

# Failure of Master Site


Suppose Site B detects the failure of master site M



**Case 1: Master site has indeed failed**



**Case 2: Partitioned network**




**Case 3: Master site has failed & Partitioned network**


# Failure of Master Site (cont.)

## Options:


- Wait for recovery of master site / network
  - ▶ Not good for availability
- Elect a new master site
  - ▶ Need to ensure at most one partition of replicas has an operational master site



Case 1



Case 2



Case 3


# Failure of Master Site (cont.)

## Electing a new master site


1. Choose a partition of replicas  $P$  to remain available
  - ▶ Simple algorithm
  - ▶ Majority consensus algorithm
  - ▶ Quorum consensus algorithm
2. If  $P$  does not contain an operational master site, elect a new master site in  $P$ 
  - ▶ Consensus algorithm

# Simple Algorithm


- The partition of replicas that contains an operational master site continues to be available



Case 1



Case 2




Case 3


- Cases 1 & 3:** System becomes unavailable
- Case 2:** System remains available with the partition containing site M operational

# Majority Consensus Algorithm


- A partition of replicas is allowed to have a master site iff the partition includes a majority of the replicas
- Majority means more than  $\lfloor \frac{N}{2} \rfloor$ , where  $N$  is the number of replicas



Case 1



Case 2




Case 3

- **Cases 1 & 2:** Majority partition remains operational after election of a new master site
- **Case 3:** System becomes unavailable



# Majority Consensus Algorithm (cont.)




Case 4: Master site is operational in minority partition

- Site M stops functioning as master site
- A new master site is elected in majority partition
- Majority partition remains operational

# Quorum Consensus Algorithm

- Each replica has a non-negative weight
- A partition of replicas is allowed to have a master site iff the total weight of the partition exceeds half the total weight of all the replicas




# CAP Theorem

- CAP
  - ▶ Data **C**onsistency
  - ▶ System **A**vailability
  - ▶ Tolerance to Network **P**artitions
- **CAP Theorem:** When there's a partitioned network, forfeit either consistency or availability
- **Forfeit consistency**
  - ▶ Resume execution on a selected partition
  - ▶ Data could become inconsistent if the selected partition requires a new master site
- **Forfeit availability**
  - ▶ Wait for network to recover before resuming execution

# Partitioned Network



Consistency OK



**Selected partition with master site M**

# Partitioned Network

Could be inconsistent



**Selected partition  
w/o master site**

# References

- T. Özsu & P. Valdureiz, *Data Replication*, Chapter 6, Principles of Distributed Database Systems, 4<sup>th</sup> Edition, 2020