
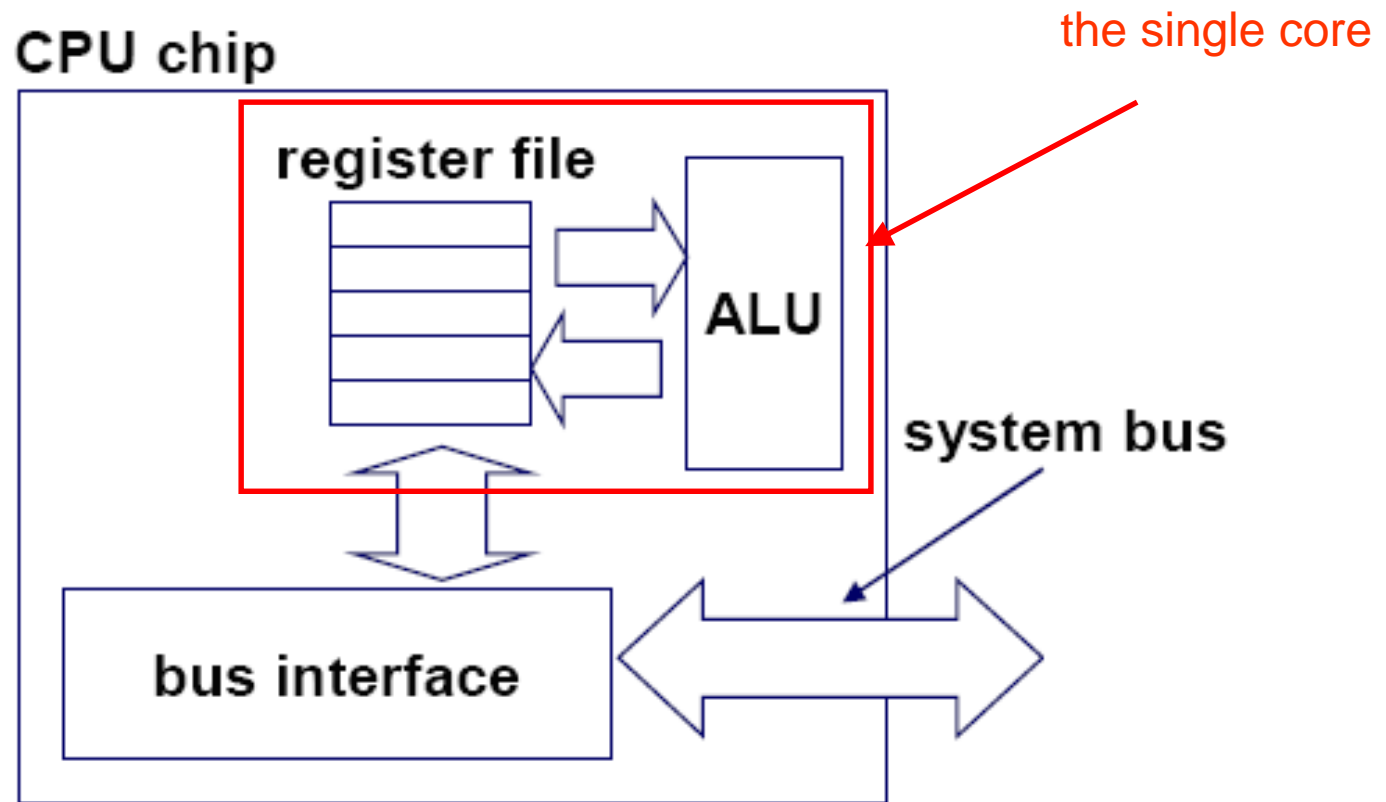


Single-core computer



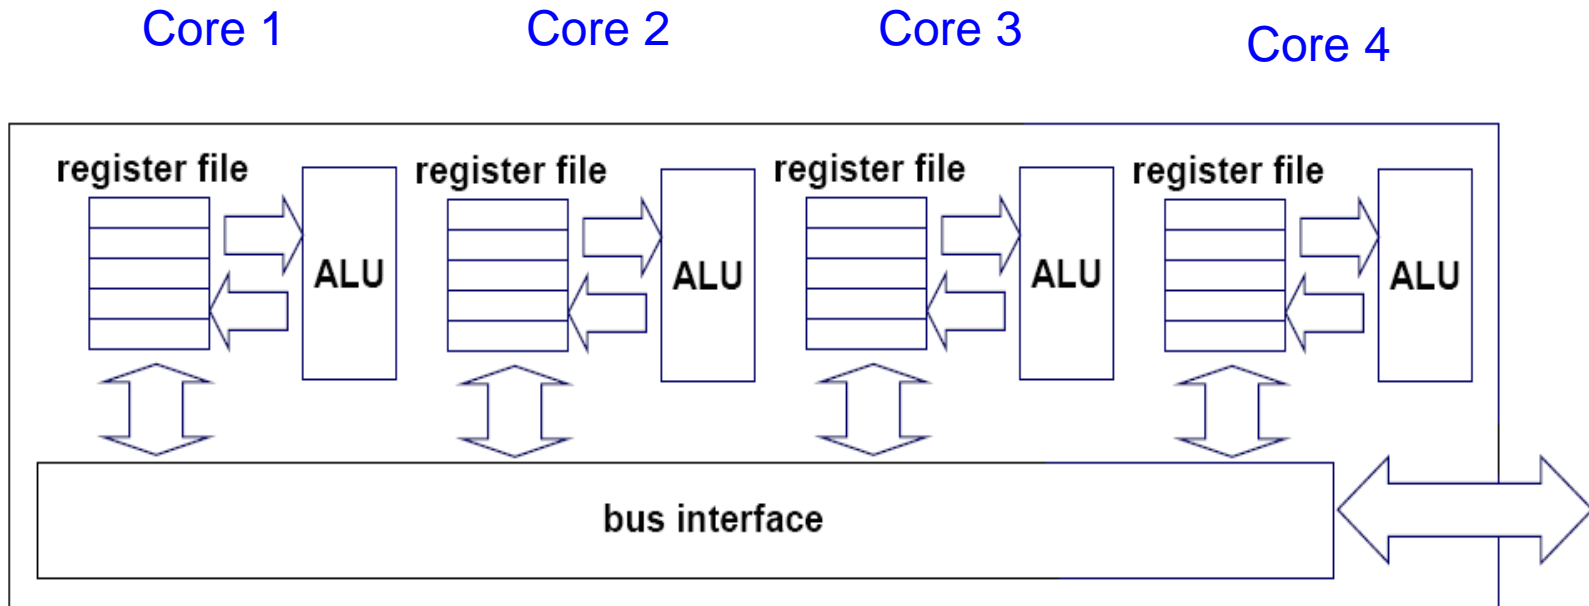
15-213, S'06

Single-core CPU chip




Multi-core architectures

- Replicate multiple processor cores on a single die.




Multi-core CPU chip


- The cores fit on a single processor socket
- Also called CMP (Chip Multi-Processor)



Cores run in parallel



Within each core, threads are time-sliced (just like on a uniprocessor)



Interaction with the Operating System

- OS perceives each core as a separate processor
- OS scheduler maps threads/processes to different cores
- Most major OS support multi-core today: Windows, Linux, Mac OS X, ...

Instruction-level parallelism

- Parallelism at the machine-instruction level
- The processor can re-order, pipeline instructions, split them into microinstructions, do aggressive branch prediction, etc.
- Instruction-level parallelism enabled rapid increases in processor speeds over the last 15 years

Thread-level parallelism (TLP)

- This is parallelism on a more **coarser scale**
- Server can serve each client in a separate thread (Web server, database server)
- A computer game can do AI, graphics, and physics in three separate threads
- Single-core superscalar processors cannot fully exploit TLP
- Multi-core architectures: explicitly exploiting TLP

Multiprocessor memory types

- Shared memory:
In this model, there is one (large) common shared memory for all processors
- Distributed memory:
In this model, each processor has its own (small) local memory, and its content is not replicated anywhere else

Multi-core processor is a special kind of a multiprocessor:


All processors are on the same chip

- Multi-core processors are **MIMD**:
Different cores execute different threads
(**M**ultiple **I**nstructions), operating on different parts of memory (**M**ultiple **D**ata).
- Multi-core is a shared memory multiprocessor:
All cores share the same memory

A technique complementary to multi-core: Simultaneous multithreading

Processor pipeline
can get stalled: *Why?*

- Waiting for the result of a long floating point (or integer) operation
- Waiting for data to arrive from memory
- Other execution units wait;

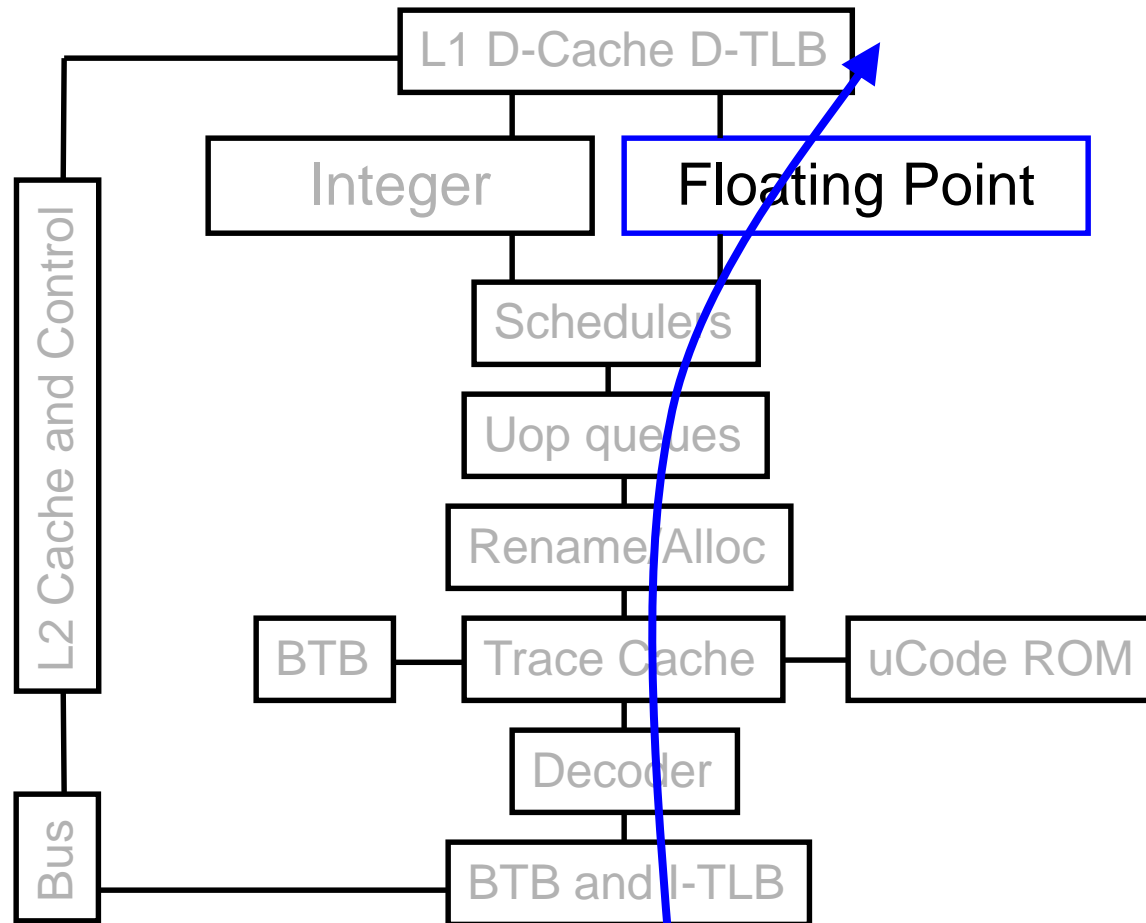


Source: Intel

Simultaneous multithreading (SMT)

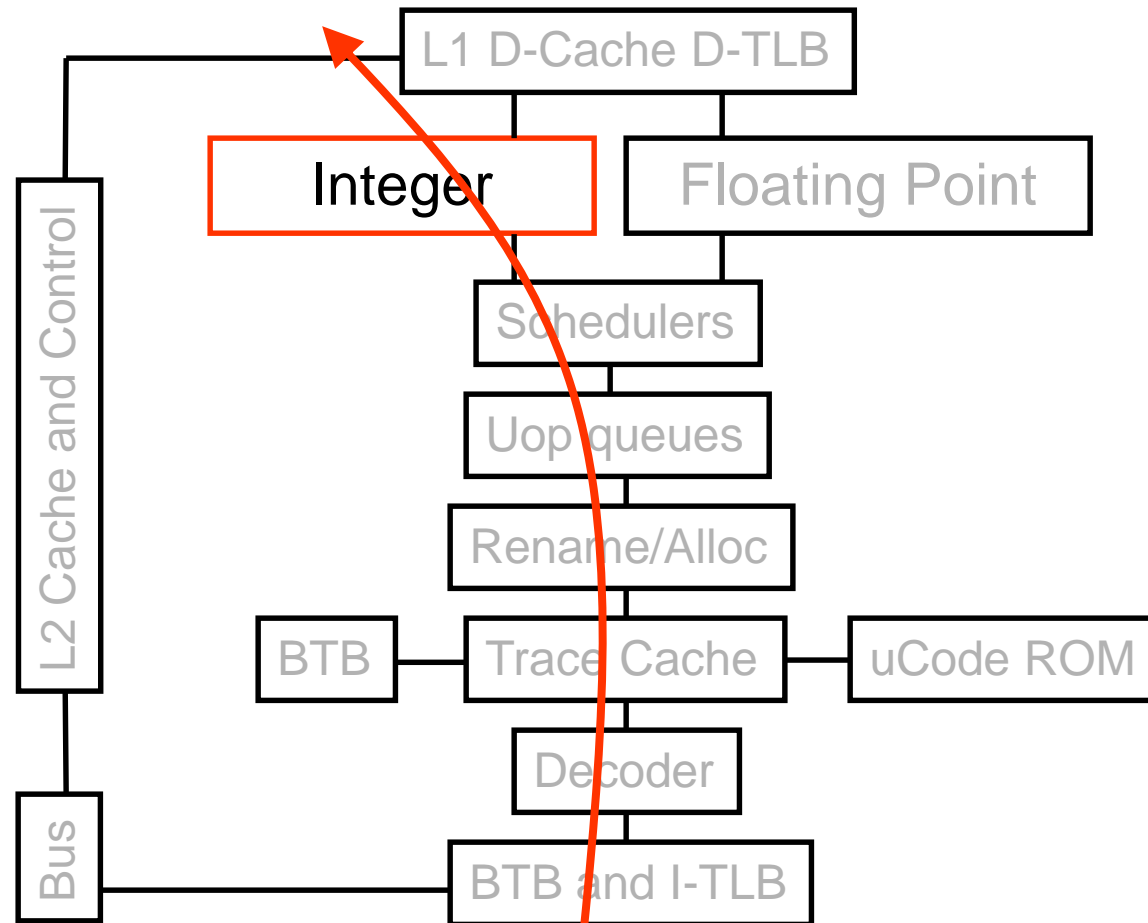
- Permits **multiple independent threads** to execute SIMULTANEOUSLY on the SAME core
- Weaving together multiple “threads” on the same core
- **Example:** if one thread is waiting for a floating point operation to complete, another thread can use the integer units

Without SMT, only a single thread can run at any given time



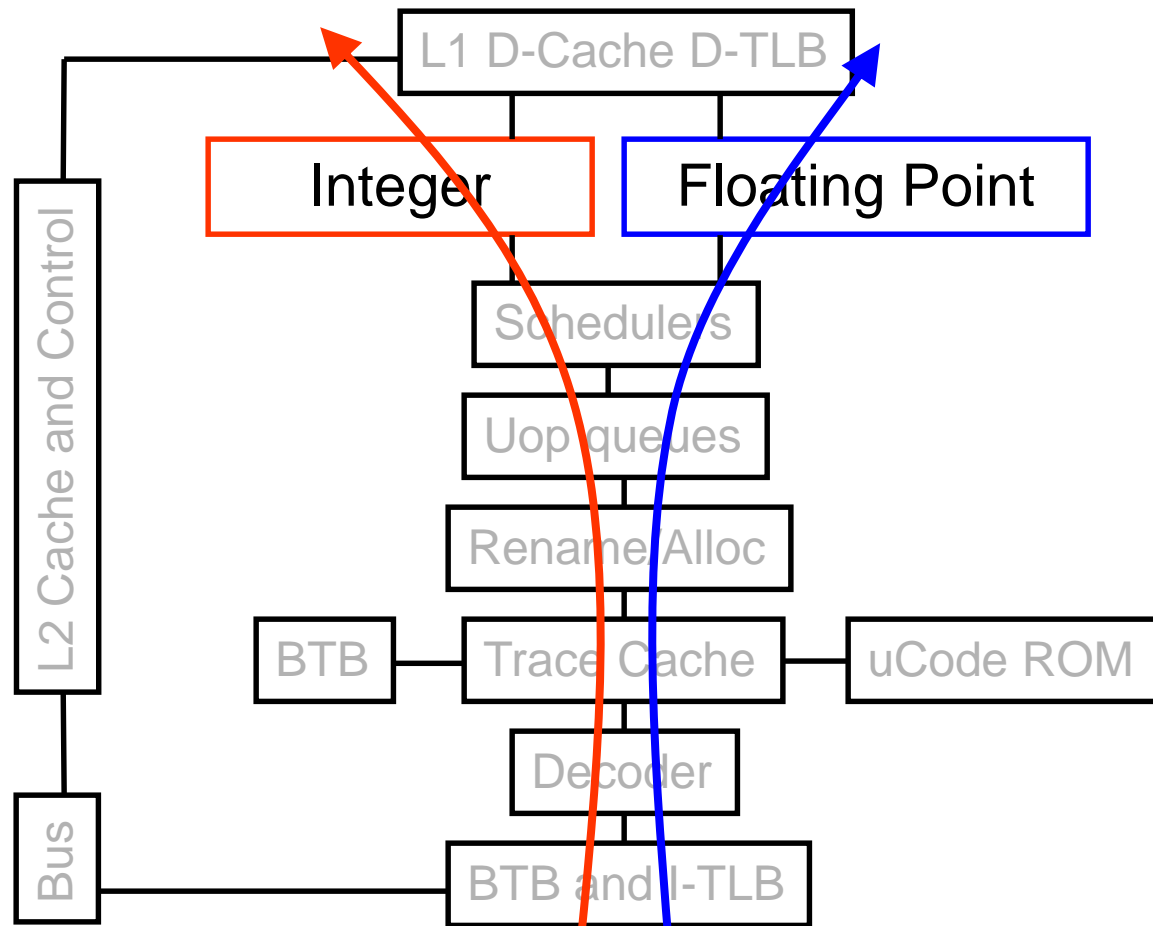
Thread 1: floating point

Without SMT, only a single thread can run at any given time



Thread 2:
integer operation

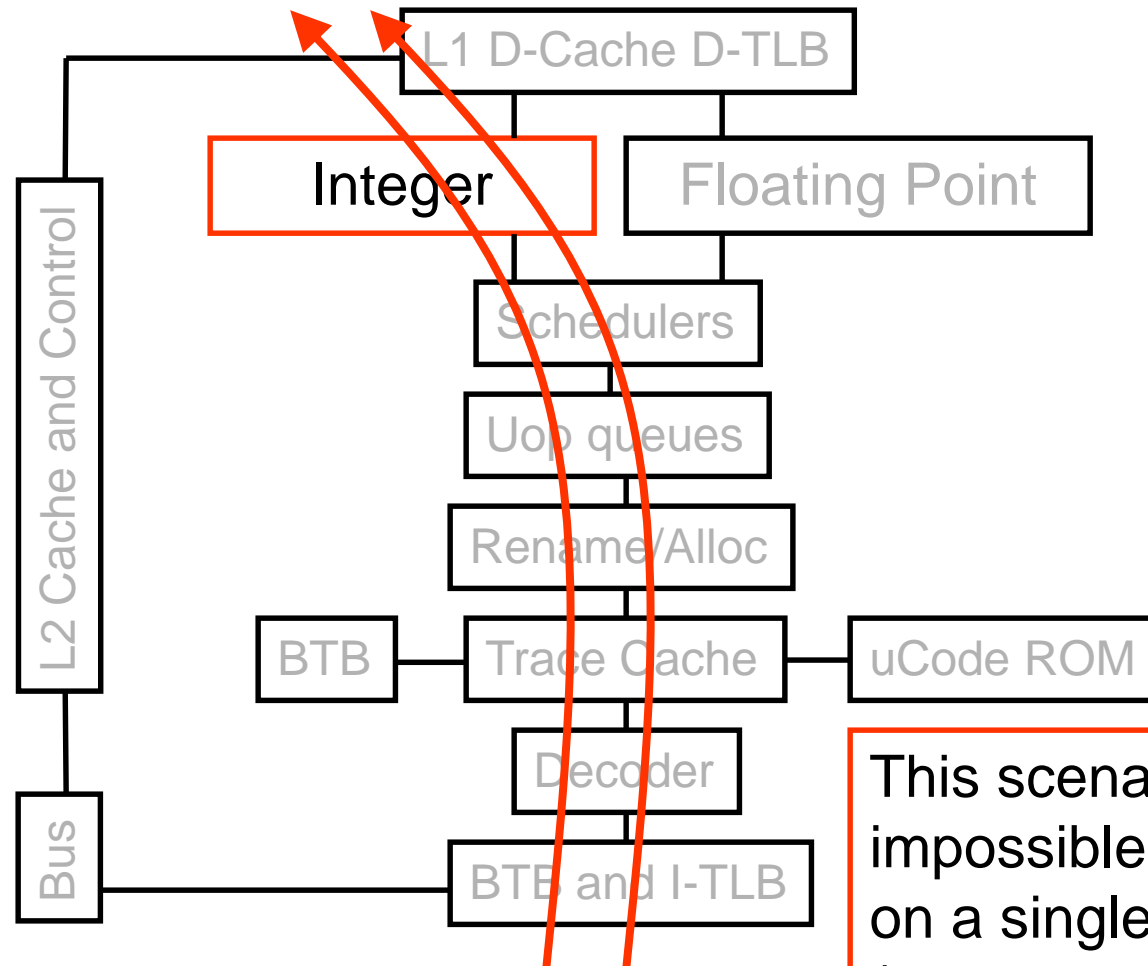
SMT processor: both threads can run concurrently



Thread 2:
integer operation

Thread 1: floating point

But can't simultaneously use the same functional unit



Thread 1 Thread 2
IMPOSSIBLE

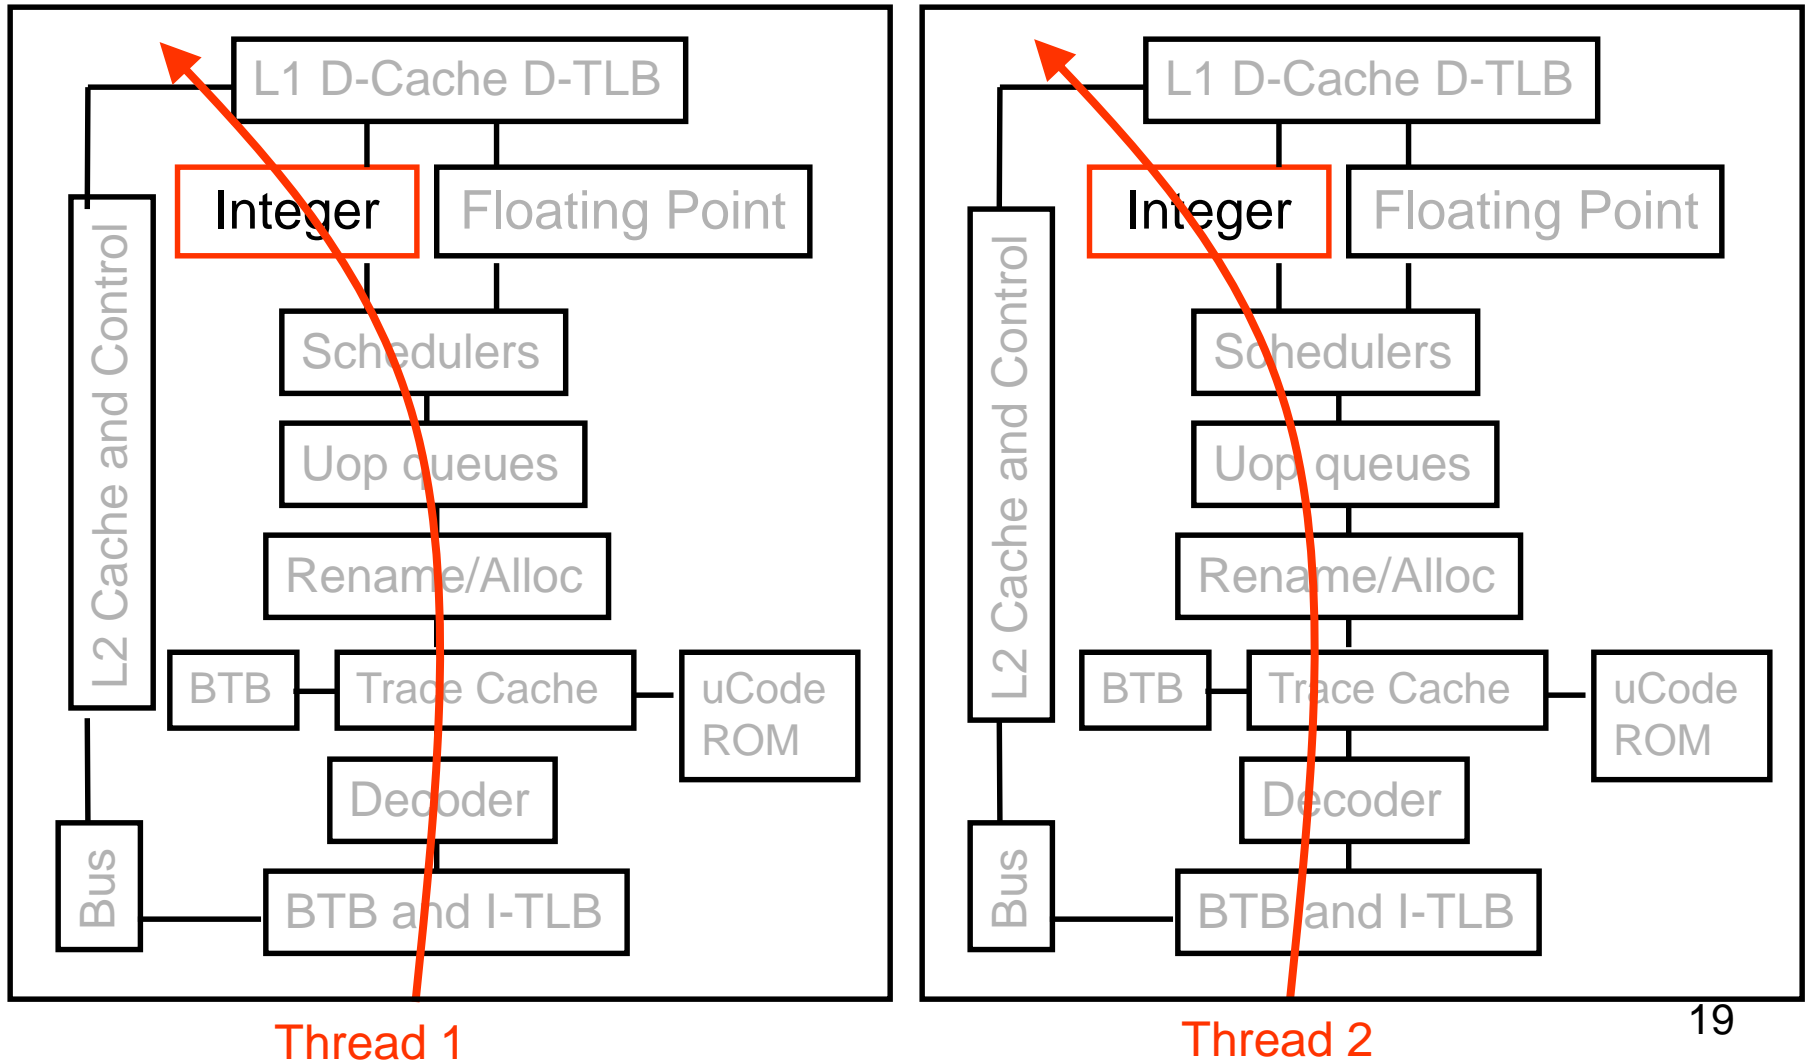
This scenario is impossible with SMT on a single core (assuming a single integer unit)

SMT not a “true” parallel processor

- Enables better threading (e.g. up to 30%)
- OS and applications perceive each simultaneous thread as a separate “virtual processor”
- The chip has only a single copy of each resource
- Compare to multi-core:
each core has its own copy of resources

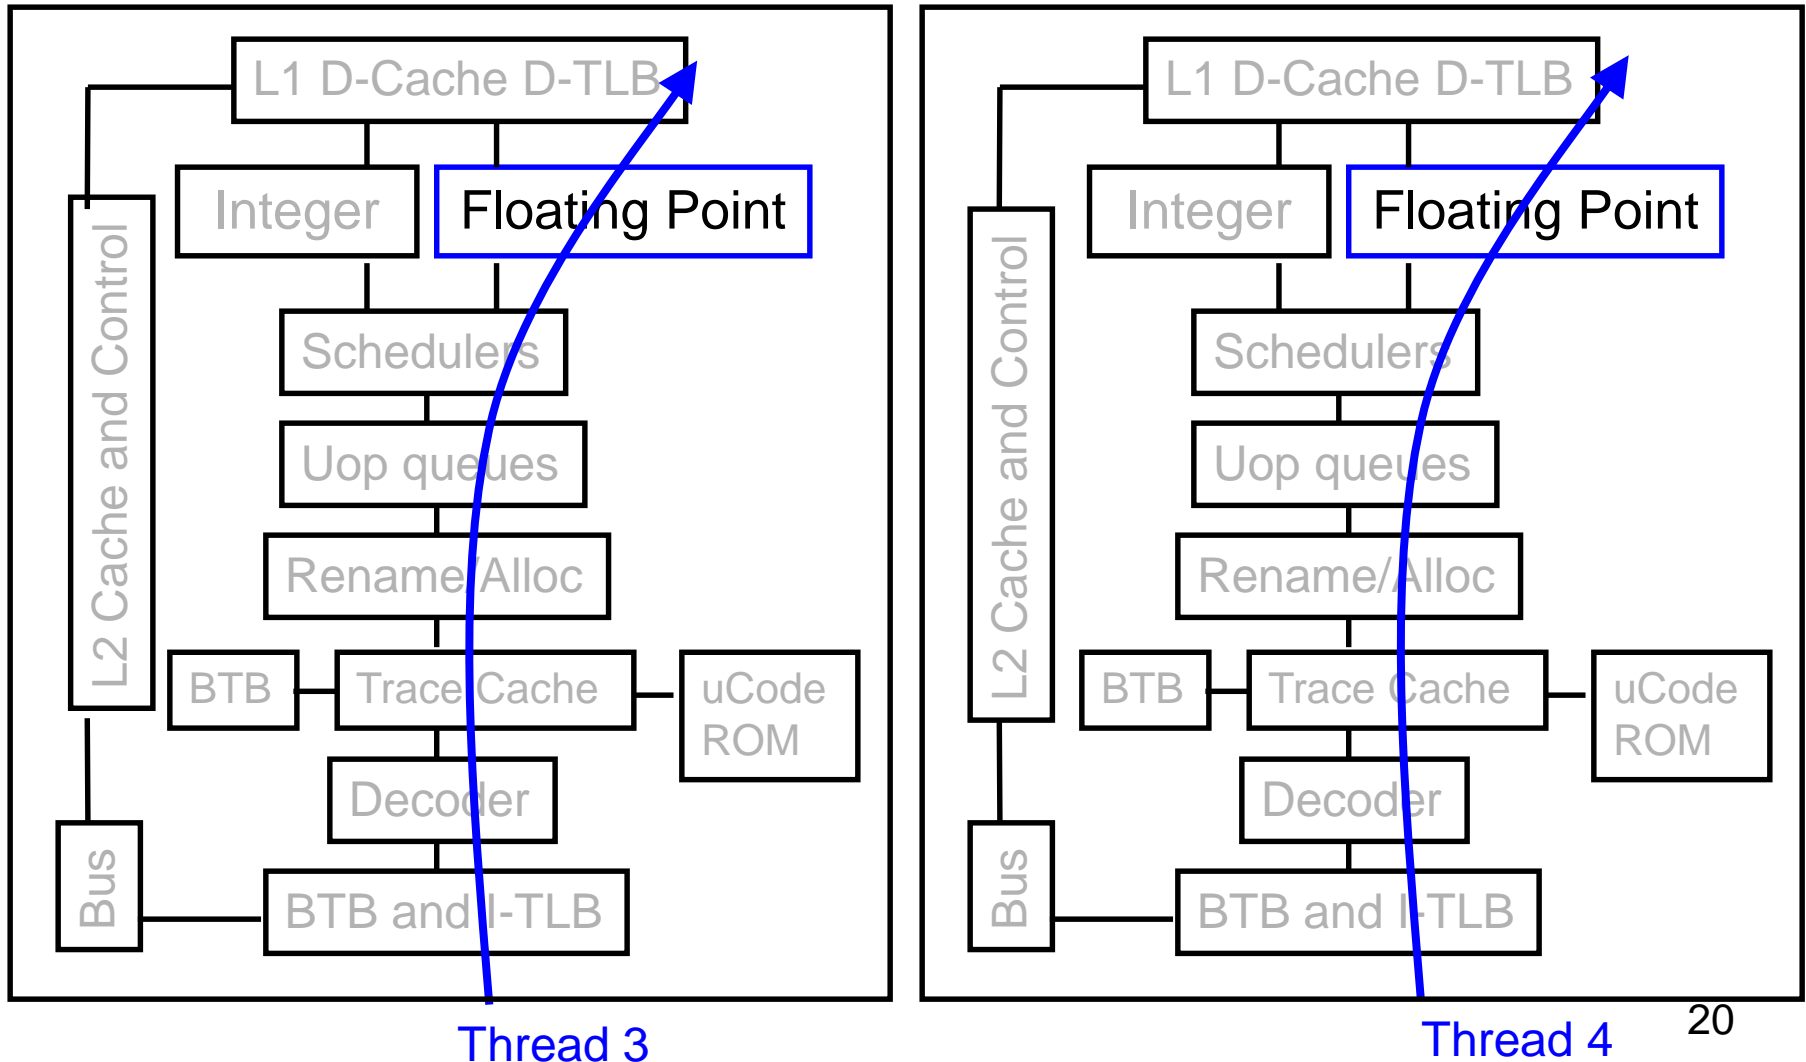
Multi-core:

threads can run on separate cores



Multi-core:

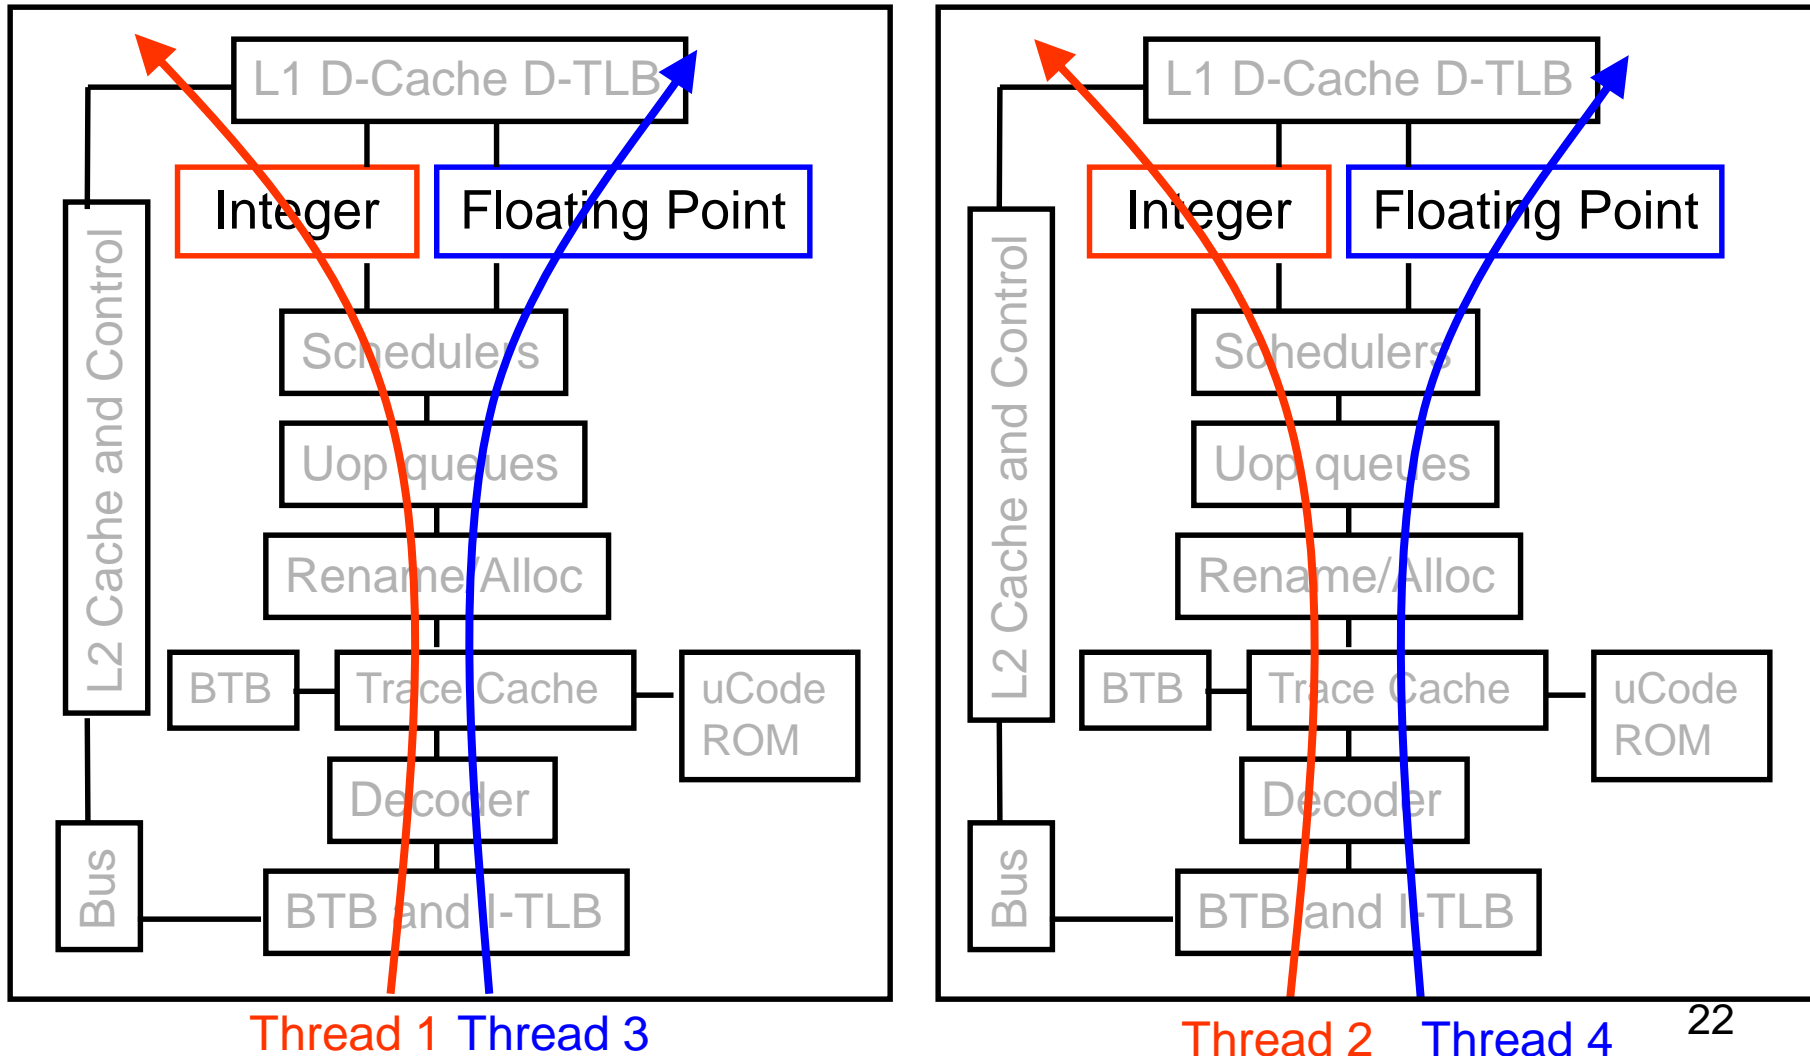
threads can run on separate cores



Combining Multi-core and SMT

- Cores can be SMT-enabled (or not)
- The different combinations:
 - Single-core, non-SMT: standard uniprocessor
 - Single-core, with SMT
 - Multi-core, non-SMT
 - Multi-core, with SMT
- The number of SMT threads:
2, 4, or sometimes 8 simultaneous threads
- Intel calls them “hyper-threads”

SMT Dual-core: all four threads can run concurrently



Comparison: multi-core vs SMT


- Multi-core:
 - Since there are several cores, each is smaller and not as powerful (but also easier to design and manufacture)
 - However, powerful with thread-level parallelism
- SMT
 - Can have one large and fast superscalar core
 - Great performance on a single thread
 - Mostly still only exploits instruction-level parallelism

The memory hierarchy


- If simultaneous multithreading only:
 - all caches are shared
- Multi-core chips:
 - L1 caches private
 - L2 caches private in some architectures and shared in others
- Memory is always shared

Typical Dual Core Features

- Dual-core Intel Xeon processors
- Each core is hyper-threaded
- Private L1 caches
- Shared L2 caches




Designs with private L2 caches



Both L1 and L2 are private

Examples: AMD Opteron,
AMD Athlon, Intel Pentium D



A design with L3 caches

Example: Intel Itanium 2 26

Private vs shared caches

- Advantages of private:
 - They are closer to core, so faster access
 - Reduces contention
- Advantages of shared:
 - Threads on different cores can share the same cache data
 - More cache space available if a single (or a few) high-performance thread runs on the system

Example: Multi-threaded way of solving a computational problem

Typical application is heat wave propagation computation:

-- Averaging and testing for convergence

- (a) Single Threaded Solution &
- (b) Multi-threaded approach using Shared Memory Multiprocessor System
- (c) Message Passing approach (suitable for HPC)

Procedure Solve(A)

/* Averaging operation 5 elements */

begin

diff = done = 0;

SINGLE THREADED EXECUTION

while (!done) do

diff = 0;

for i \leftarrow 1 to n do

for j \leftarrow 1 to n do

temp = A[i,j]; /* saving the original value */

A[i,j] \leftarrow 0.2 * (A[i,j] + neighbors);

diff += abs(A[i,j] - temp); /* difference computation */

end for

end for

if (diff < TOL) then done = 1; /* convergence testing */

end while

end procedure

```

int n, nprocs;
float **A, diff;
LOCKDEC(diff_lock);
BARDEC(bar1);

```

```

main()
begin
  read(n); read(nprocs);
  A ← G_MALLOC();
  initialize (A);
  CREATE (nprocs, Solve, A);
  WAIT_FOR_END (nprocs);
end main

```

MULTI-THREDED EXECUTION - SM

```

procedure Solve(A)
  int i, j, pid, done=0;
  float temp, mydiff=0;
  int mymin = 1 + (pid * n/nprocs);
  int mymax = mymin + n/nprocs - 1;
  while (!done) do
    mydiff = diff = 0;
    BARRIER(bar1, nprocs);
    for i ← mymin to mymax
      for j ← 1 to n do
        ...
      endfor
    endfor
    LOCK(diff_lock);
    diff += mydiff;
    UNLOCK(diff_lock);
    BARRIER (bar1, nprocs);
    if (diff < TOL) then done = 1;
    BARRIER (bar1, nprocs);
  endwhile

```

```

main()
  read(n); read(nprocs);
  CREATE (nprocs-1, Solve);
  Solve();
  WAIT_FOR_END (nprocs-1);

procedure Solve()
  int i, j, pid, nn = n/nprocs, done=0;
  float temp, tempdiff, mydiff = 0;
  myA ← malloc(...)
  initialize(myA);
  while (!done) do
    mydiff = 0;
    if (pid != 0)
      SEND(&myA[1,0], n, pid-1, ROW);
    if (pid != nprocs-1)
      SEND(&myA[nn,0], n, pid+1, ROW);
    if (pid != 0)
      RECEIVE(&myA[0,0], n, pid-1, ROW);
    if (pid != nprocs-1)
      RECEIVE(&myA[nn+1,0], n, pid+1, ROW);

```

```

    for i ← 1 to nn do
      for j ← 1 to n do
        ...
      endfor
    endfor
  if (pid != 0)
    SEND(mydiff, 1, 0, DIFF);
    RECEIVE(done, 1, 0, DONE);
  else
    for i ← 1 to nprocs-1 do
      RECEIVE(tempdiff, 1, *, DIFF);
      mydiff += tempdiff;
    endfor
    if (mydiff < TOL) done = 1;
    for i ← 1 to nprocs-1 do
      SEND(done, 1, i, DONE);
    endfor
  endif
endwhile

```

MESSAGE PASSING APPROACH

Sizes of Cache for I & D? Empirical study
is possible....*Stay tuned!*

*Let us now look into three types of Multithreaded
Processors.*

Acknowledgements: Thanks to Dr. Jernej Barbic, 2007